

Discussion #5

Diwanshu, Elena

July 8, 2020

Logistics

- HW3 due 7/12, 11:59 PM.
- Checkpoint 3 due 7/15, 11:59 PM.

Outline

- B+ Tree → Insertion, Deletion.
- Buffer Manager → LRU, MRU

Why B+ Tree?

→ Most queries can be executed more quickly if the values are stored in order.


→ But it's not practical to hope to store all the rows in the table one after another, in sorted order, because this requires rewriting the entire table with each insertion or deletion of a row.

Requirements of B+ Tree

- Two types of nodes: index (internal) nodes and data (leaf) nodes. Each node is one disk page.
- Each non-leaf (“internal”) node has $d \leq m \leq 2d$ entries
 - Minimum **50%** occupancy
- Root node has $1 \leq m \leq 2d$ entries
- We call **d** the **order** of the tree

B+ Tree


- The n entries in a node define $n+1$ ranges.



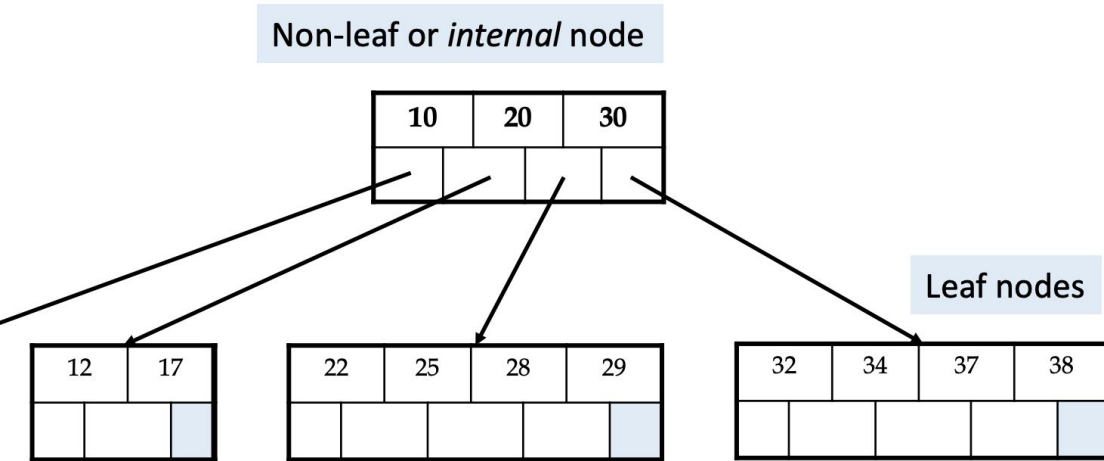
B+ Tree

- For each range, in a non-leaf node, there is a pointer to another node with entries in that range.

Non-leaf or *internal* node

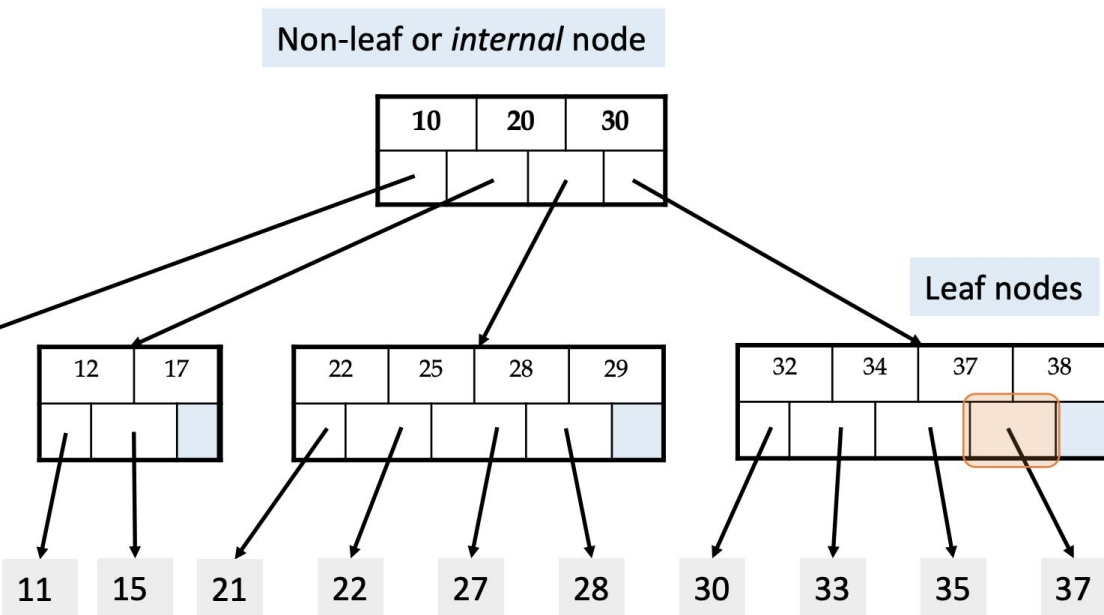


B+ Tree




- Leaf nodes also have between **d** and **2d** entries, and are different in that:

B+ Tree



- Leaf nodes also have between **d** and **2d** entries, and are different in that:
 1. Their entry slots contain pointers to data records.

B+ Tree

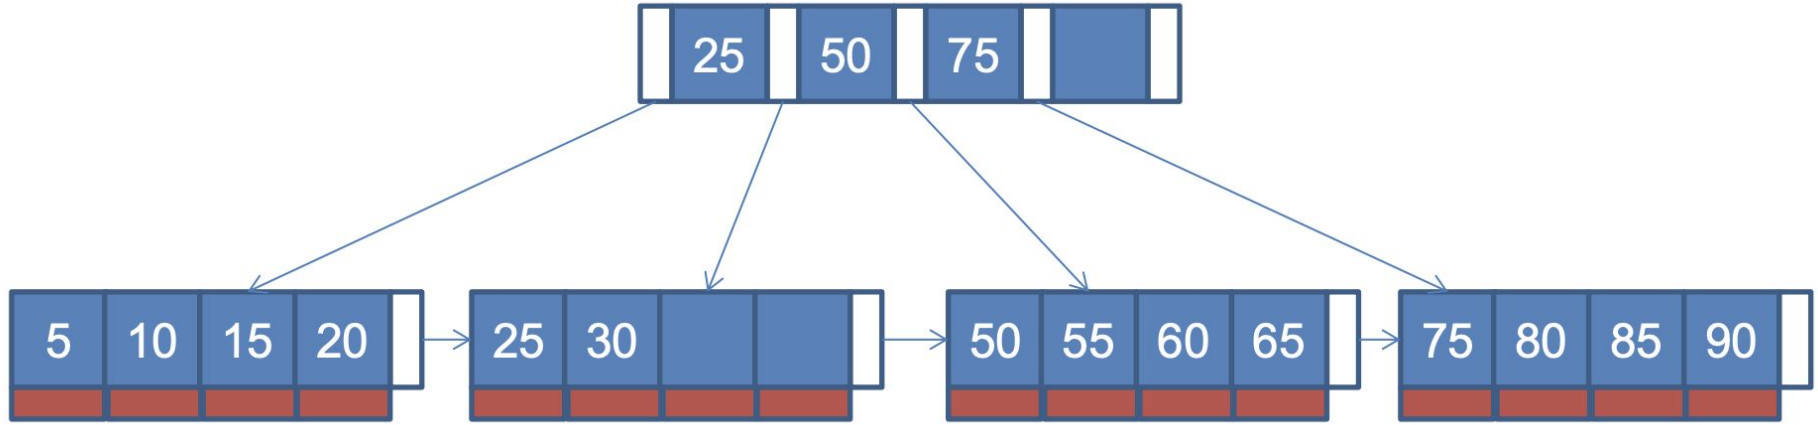


- Leaf nodes also have between **d** and **2d** entries, and are different in that:
 1. Their entry slots contain pointers to data records.
 2. They contain a pointer to the next leaf node as well, for faster sequential traversal.

Inserting

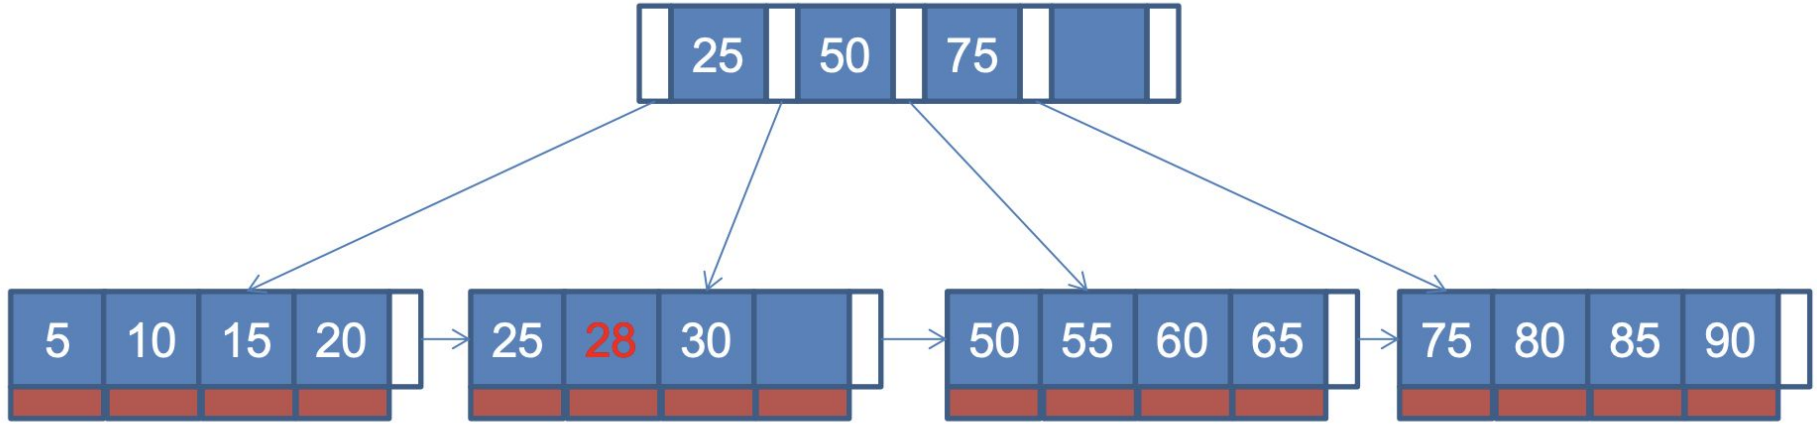
- Find leaf where new key **x** belongs
- If leaf is not full then add **x** to leaf
- If leaf is full then split leaf and push splitting key up to parent
 - If parent is full then split parent and push splitting key up
 - If we need to add a key to the root and the root is full then split the root and create a new parent as root

Inserting



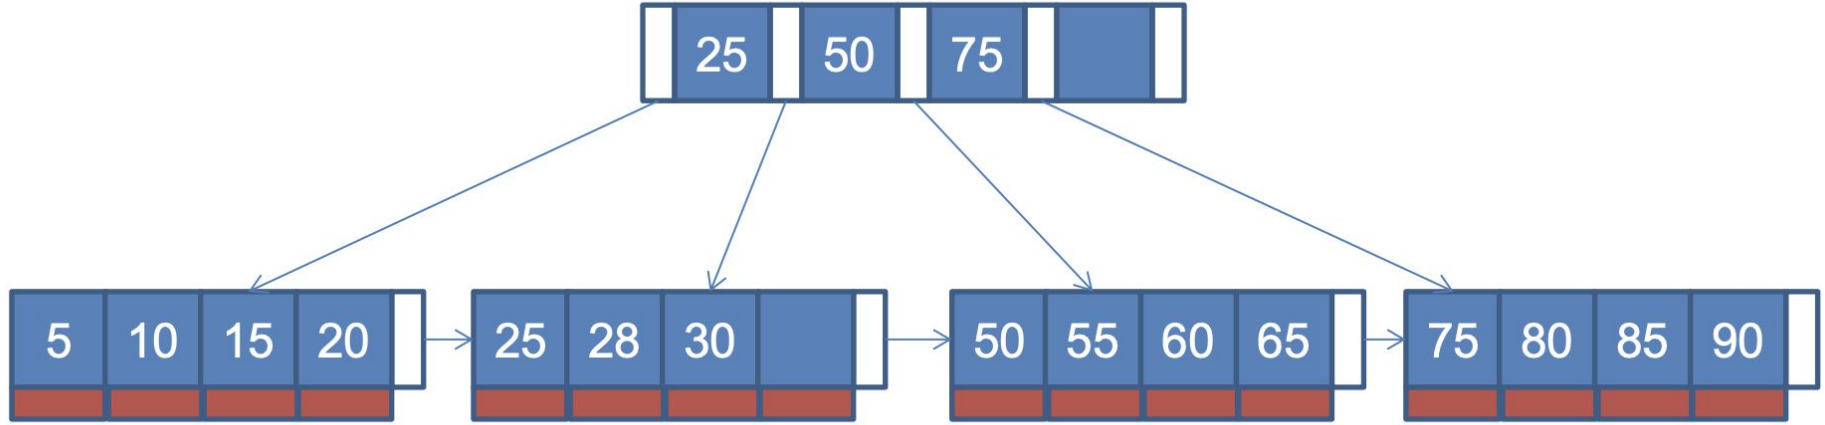
- Insert 28
 - Leaf is not full so we can just add it

Inserting



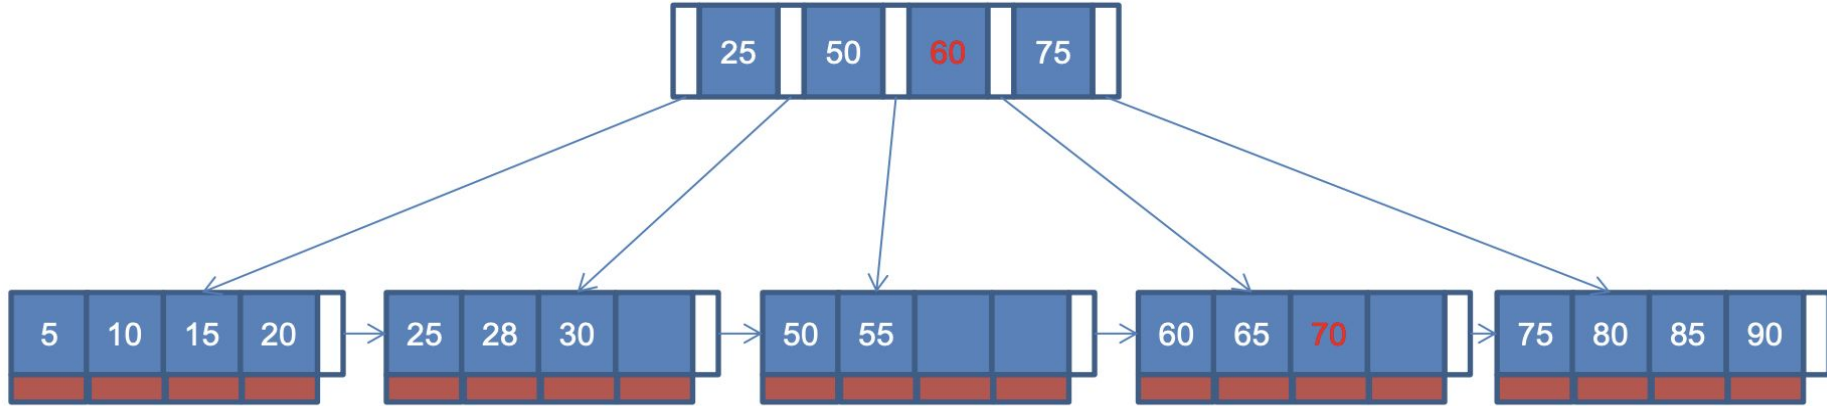
- Insert 28
 - Leaf is not full so we can just add it

Inserting



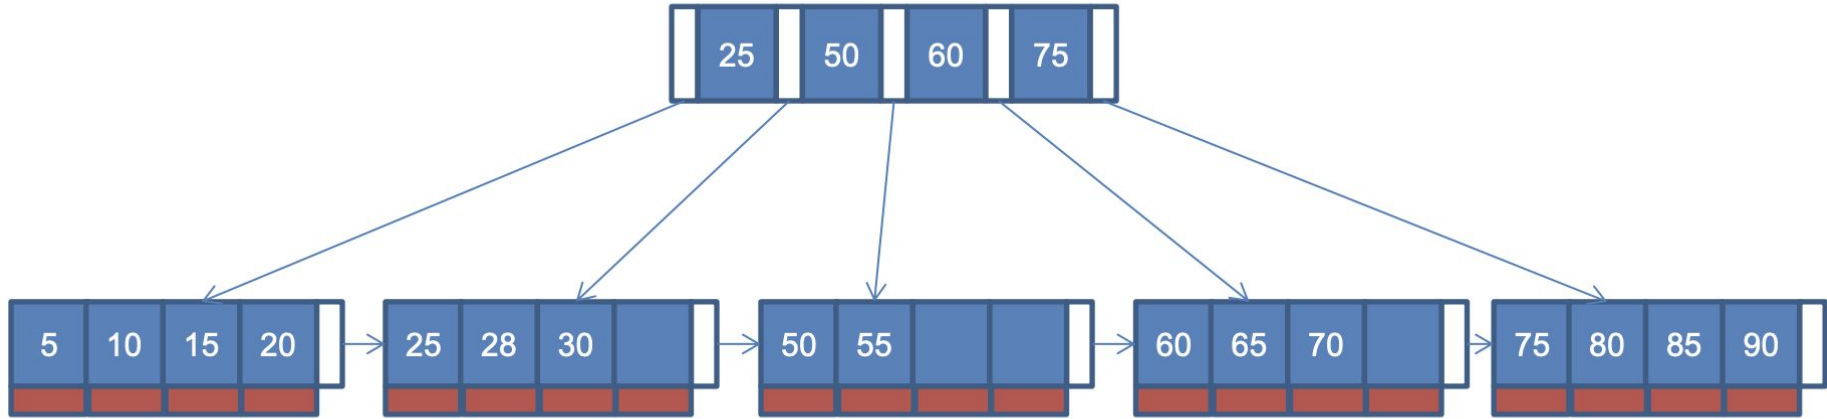
- Insert 70
 - Should go in leaf with (50, 55, 60, 65), but leaf is full → Need to split the leaf and then insert.

Inserting



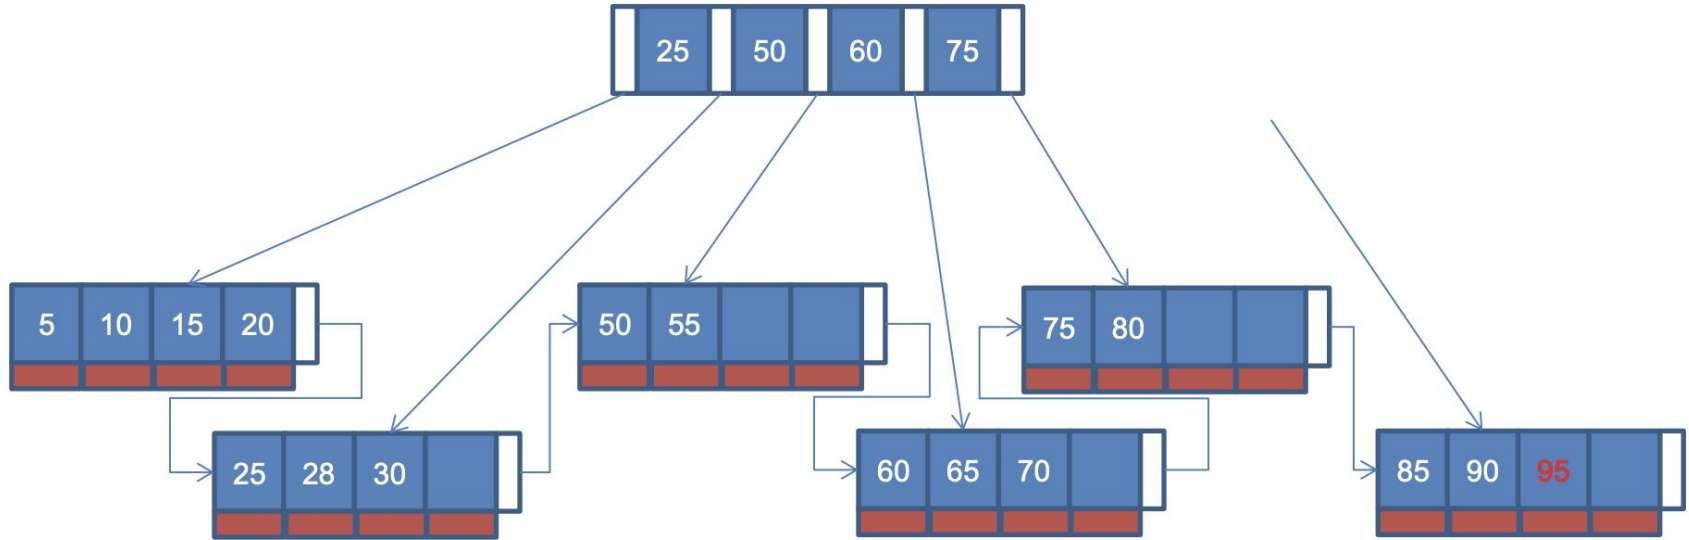
- Insert 70
 - Should go in leaf with (50, 55, 60, 65), but leaf is full → Need to split the leaf and then insert → **60 is pushed up to the parent**

Inserting



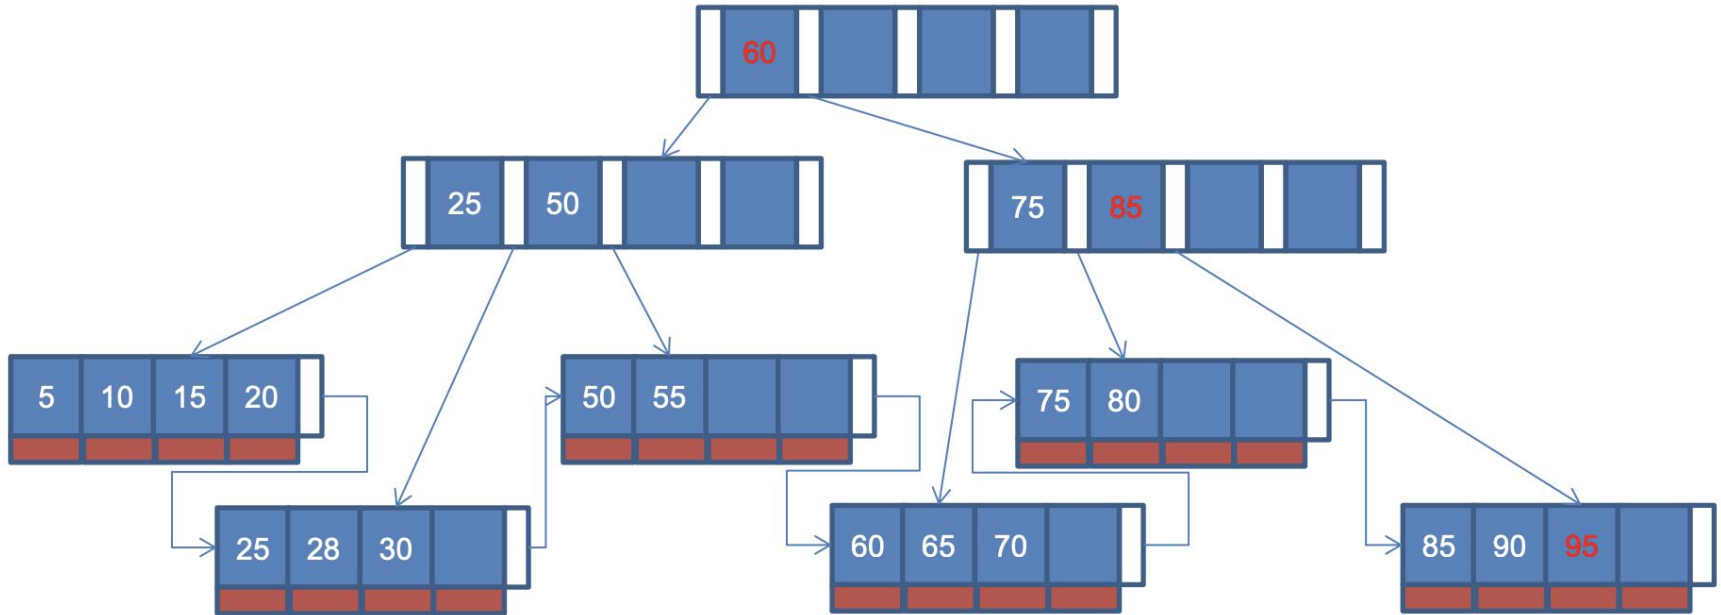
- Insert 95
 - Should go in leaf with (75, 80, 85, 90), but leaf is full → Need to split the leaf and then insert
 - 85 is pushed up to the parent, but parent is full! → Need to split parent node then insert 85

Inserting



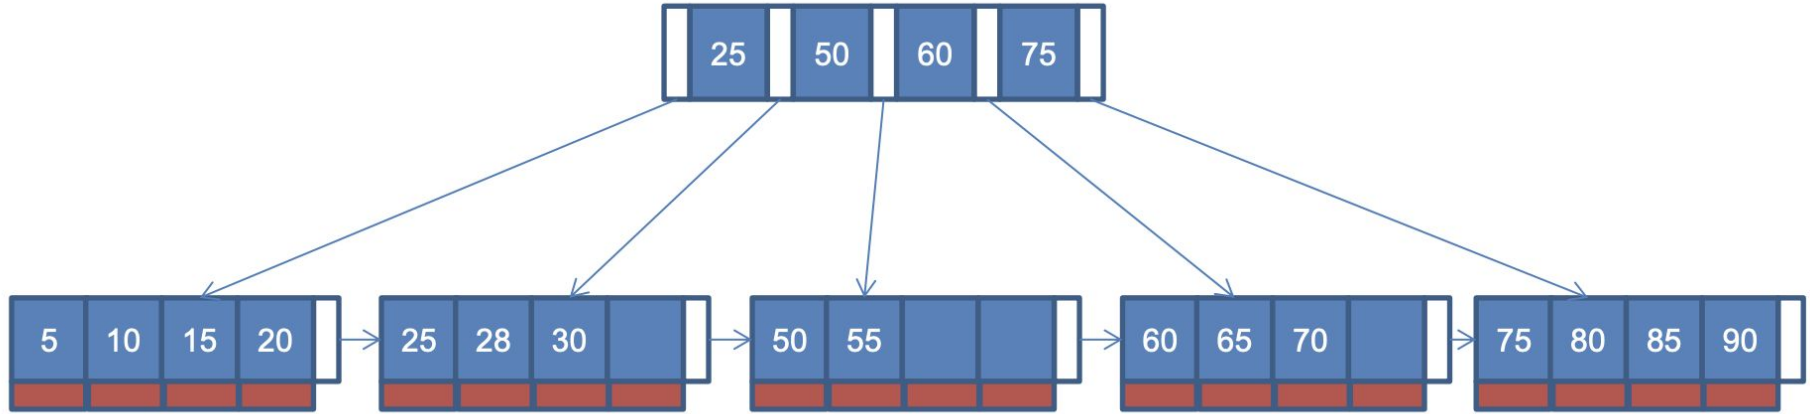
- Insert 95
 - Should go in leaf with (75, 80, 85, 90), but leaf is full → Need to split the leaf and then insert
 - 85 is pushed up to the parent, but parent is full! → Need to split parent node then insert 85

Inserting



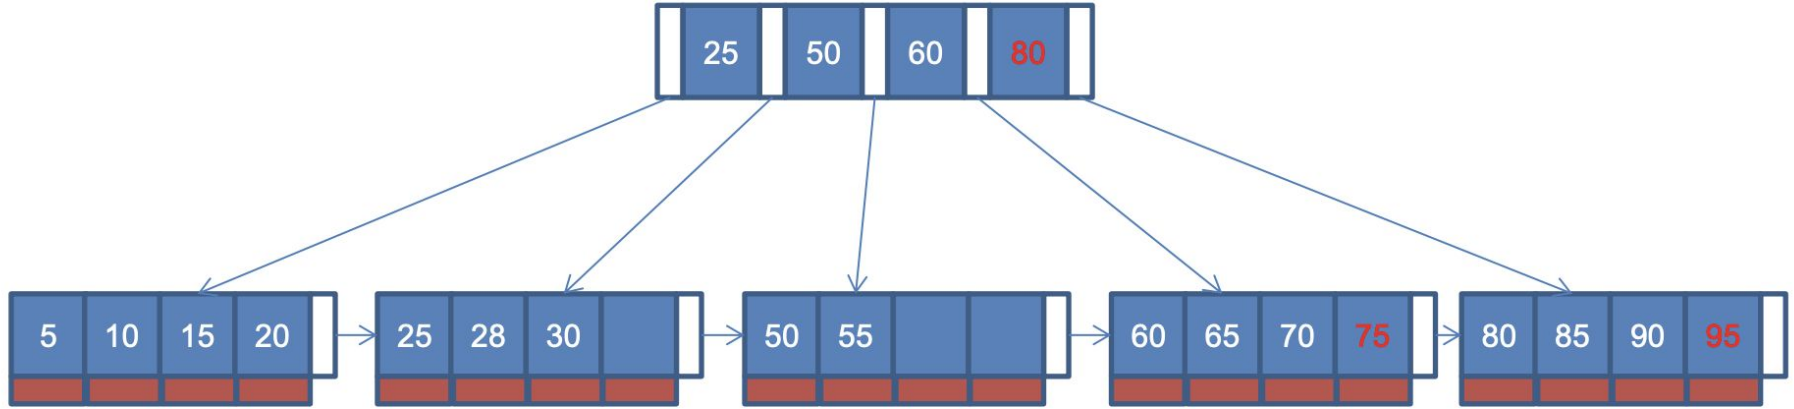
- Insert 95
 - Should go in leaf with (75, 80, 85, 90), but leaf is full → Need to split the leaf and then insert
 - 85 is pushed up to the parent, but parent is full! → Need to split parent node then insert 85
 - **60 is then pushed up to the new root**

Inserting Alternative



- If inserting into a full node with a non-full sibling, we could shift keys from the full node to the non-full sibling
 - Requires modifying parent keys

Inserting Alternative



- Insert 95
 - Should go in leaf with (75, 80, 85, 90), but leaf is full → Shift 75 to left sibling
 - Modify key in parent node → Insert 95 in newly opened node

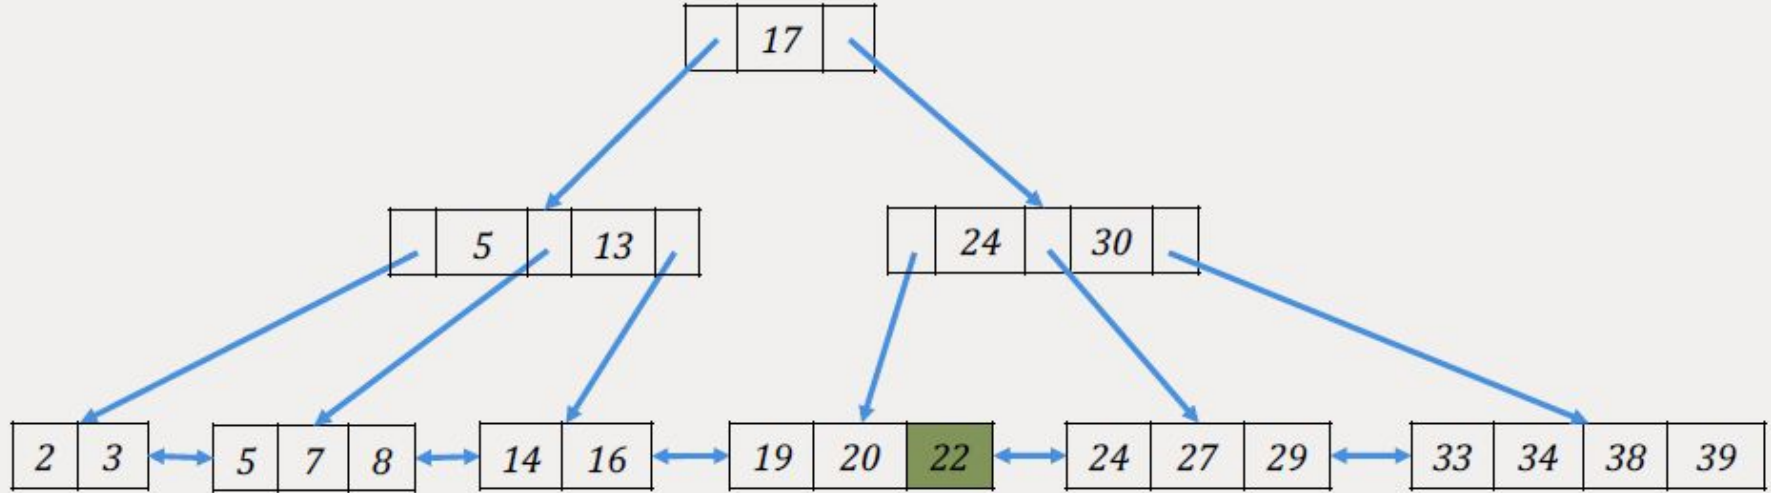
Deletion

- Find leaf node **L** where entry belongs
- Remove the entry
 - If **L** is at least half-full, DONE
 - If **L** has only $d-1$ entries
 - Try to **re-distribute**, borrowing from sibling
 - If re-distribution fails, **merge L** and sibling
- If a merge occurred, we must delete an entry from the parent of **L**

Example 1

order $d = 2$

delete 22




since by deleting 22 the node remains half-full, we simply remove it

Example 1

order $d = 2$


delete 22



Example 2

order $d = 2$

delete 20



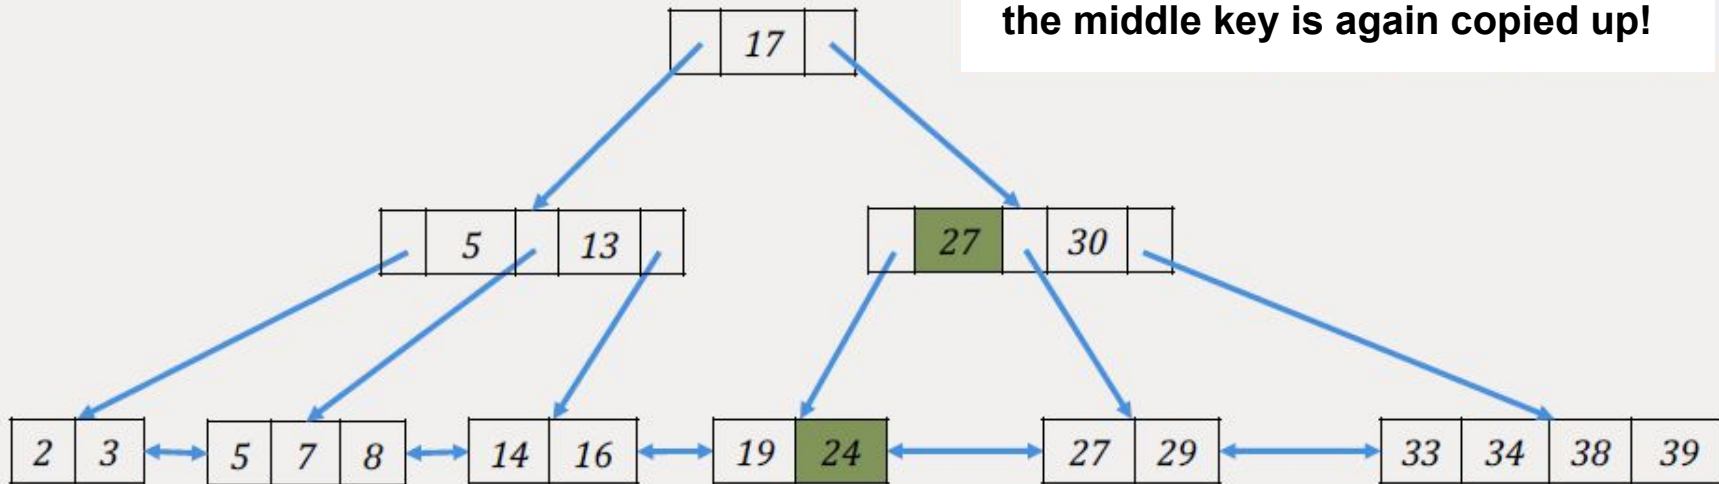
by removing 20 the node is not half-full anymore, so we attempt to redistribute!

Example 2

order $d = 2$

delete 20

the middle key is again copied up!

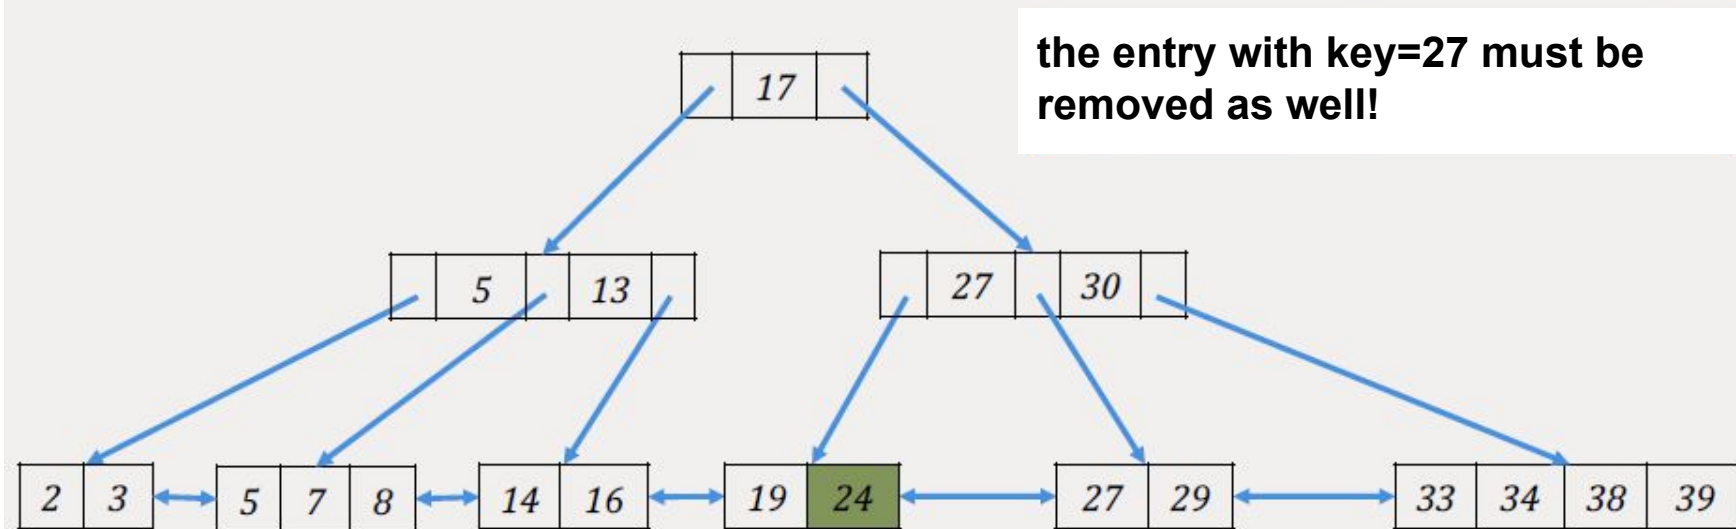


by removing 20 the node is not half-full anymore, so we attempt to redistribute!

Example 3

order $d = 2$

delete 24




in this case we have to merge nodes!

Example 3

order $d = 2$

delete 24


we are not done, since the resulting non-leaf node is not half-full!



Example 3

order $d = 2$

delete 24



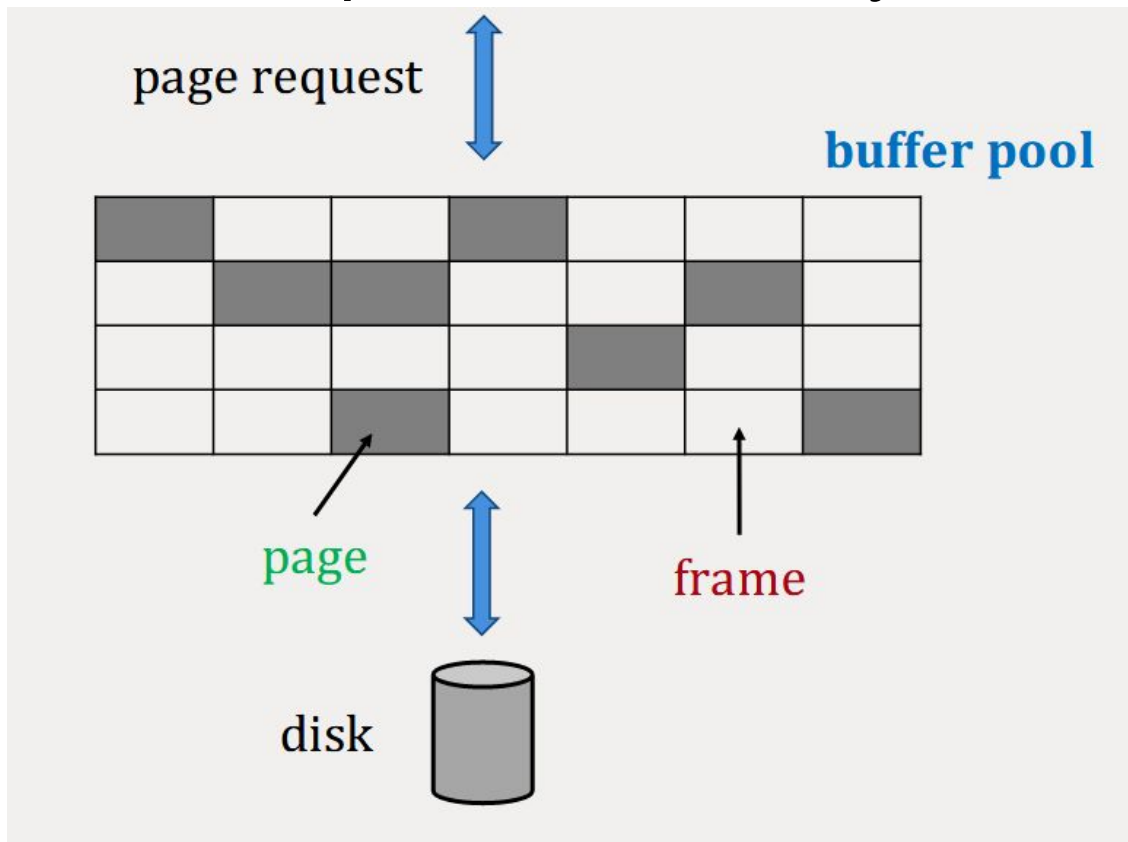
Buffer Manager

- Data must be in RAM for the DBMS to operate on it
- All the pages may not fit into the main memory

Buffer Manager: responsible for bringing pages from disk to main memory as needed

- pages brought into memory are in the **buffer pool**
- the buffer pool is partitioned into **frames**: slots for holding disk pages
- page same size with frame

Buffer Replacement Policy



How to choose a frame for replacement?

- LRU (Least Recently Used)
- MRU (Most Recently Used)
- Clock
- FIFO
- Random
- ...

LRU

- Uses a queue of pointers to frames that have **pin_count = 0**
- A page request uses frames only from the head of the queue
- When a the **pin_count** of a frame goes to 0, it is added to the end of the queue
- MRU works in a similar way

LRU - Least Recently Used

A

B

C

Page read sequence - A, B, C, A, B, C, A, B, C (*Sequential Flooding*)

Q. Miss rate?

Time	Frame 1	Frame 2
T0(A)	A	
T1(B)	A	B
T2(C)	C	B
T3(A)	C	A
T4(B)	B	A
T5(C)	B	C
T6(A)	A	C
T7(B)	A	B
T8(C)	C	B

LRU - Least Recently Used

A

B

C

Page read sequence - A, A, B, B, C, C, A, B, C

Q. Miss rate?

Time	Frame 1	Frame 2
T0(A)	A	
T1(A)	A	
T2(B)	A	B
T3(B)	A	B
T4(C)	C	B
T5(C)	C	B
T6(A)	C	A
T7(B)	B	A
T8(C)	B	C

MRU - Most Recently Used

A

B

C

Page read sequence - A, B, C, A, B, C, A, B, C (*Sequential Flooding*)

Q. Miss rate?

Time	Frame 1	Frame 2
T0(A)	A	
T1(B)	A	B
T2(C)	A	C
T3(A)	A	C
T4(B)	B	C
T5(C)	B	C
T6(A)	B	A
T7(B)	B	A
T8(C)	C	A

MRU - Most Recently Used

A

B

C

Page read sequence - A, A, B, B, C, C, A, B, C

Q. Miss rate?

Time	Frame 1	Frame 2
T0(A)	A	
T1(A)	A	
T2(B)	A	B
T3(B)	A	B
T4(C)	A	C
T5(C)	A	C
T6(A)	A	C
T7(B)	B	C
T8(C)	B	C

Thanks!