

# Discussion #6

Diwanshu, Elena

July 15, 2020

# Logistics

- Checkpoint 3 due EOD.
- Checkpoint 4 due 7/19, 11:59 PM.
- Final Project Presentation: Sign up [here](#)

# Outline

- Dynamic Hashing: Extendible Hashing
- External Sorting

# Why Dynamic Hashing?

- The problem with static hashing is that it does not expand or shrink dynamically as the size of the database grows or shrinks.
- Dynamic hashing provides a mechanism in which data buckets are added and removed dynamically and on-demand.

# Dynamic Hashing: Extendible Hashing

→ **Extendible Hashing** is a dynamic hashing method wherein directories, and buckets are used to hash data. It is an aggressively flexible method in which the hash function also experiences dynamic changes.

# Main features

- **Directories:** The directories store addresses of the buckets in pointers. An id is assigned to each directory which may change each time when Directory Expansion takes place.
- **Buckets:** The buckets are used to hash the actual data.
  - *A bucket may contain more than one pointers to it.*

# Main features

- **Global depth** of directory
  - Max # of bits needed to tell which bucket an entry belongs to
- **Local depth** of a bucket
  - # of bits used to determine if an entry belongs to this bucket

**Q. Compare:**

**Local depth ? Global depth**

# Main features

- **Global depth** of directory
  - Max # of bits needed to tell which bucket an entry belongs to
- **Local depth** of a bucket
  - # of bits used to determine if an entry belongs to this bucket

**Q. Compare:**

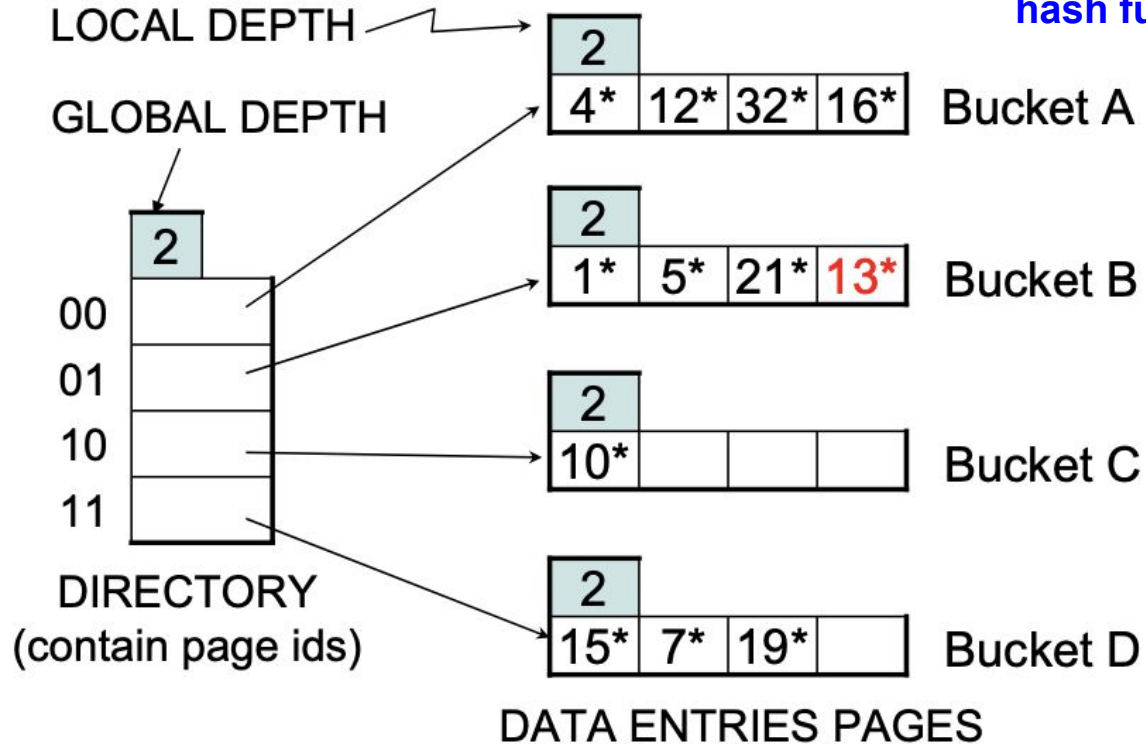
**Local depth  $\leq$  Global depth**



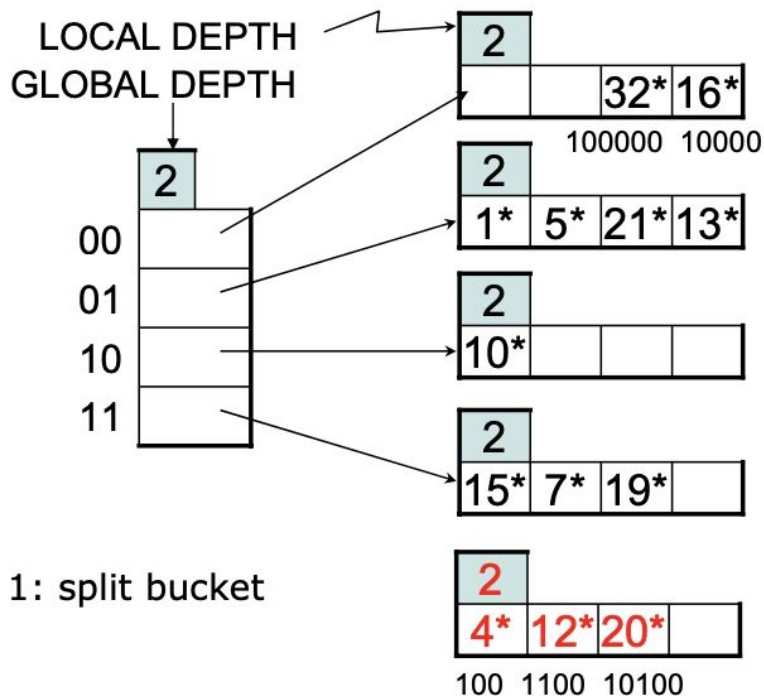
# Example

Insert entry x with  $h(x) = 13$

hash function:  $h(x) = x \bmod 4$



# Example




Insert entry  $x$  with  $h(x) = 20$

hash function:  $h(x) = x \bmod 8$

## OVERFLOW?

- If **local depth = global depth**
  - We need to double the size of table, add new bucket, redistribute data, update pointers
- If **local depth < global depth**
  - Add new bucket, redistribute data, update pointers

# Example



Step 2: double directory and increase depth

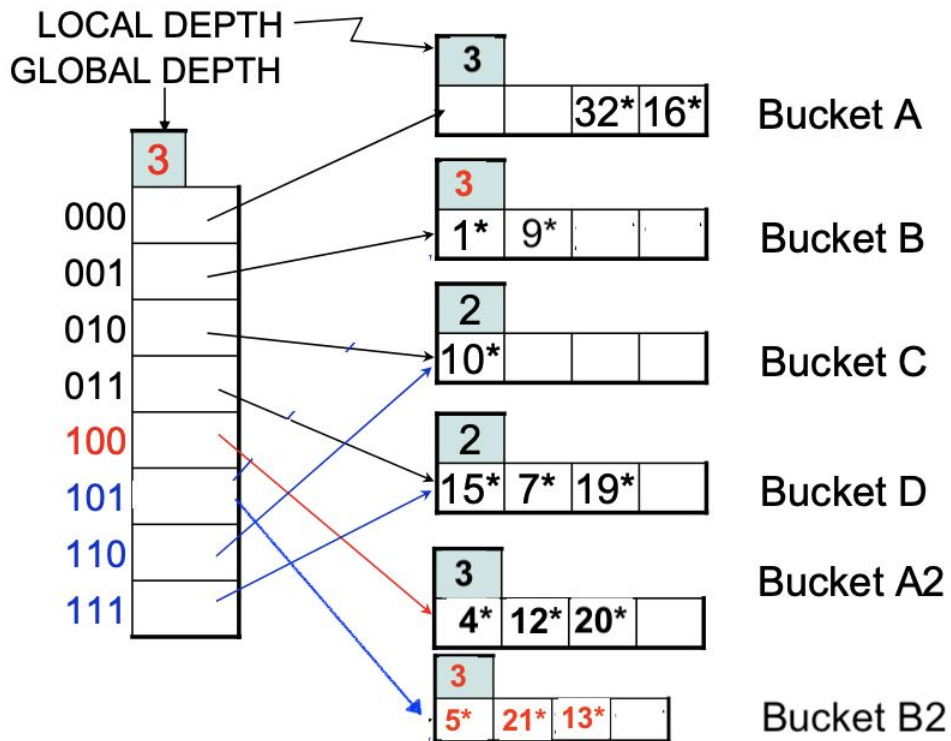
Insert entry  $x$  with  $h(x) = 20$

hash function:  $h(x) = x \bmod 8$

## OVERFLOW?

- If local depth = global depth
  - We need to double the size of table, add new bucket, redistribute data, update pointers
- If local depth < global depth
  - Add new bucket, redistribute data, update pointers

# Example



Insert entry  $x$  with  $h(x) = 9$

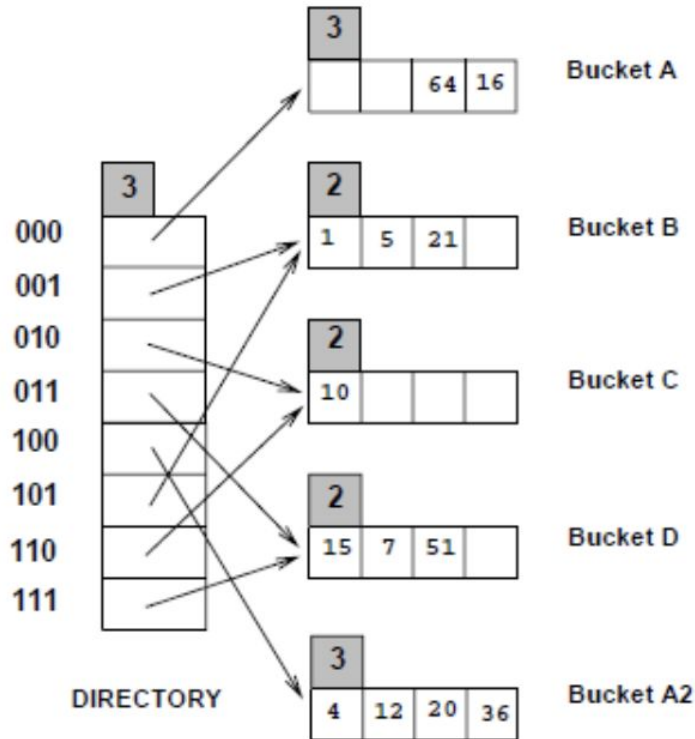
hash function:  $h(x) = x \bmod 8$

## OVERFLOW?


- If **local depth = global depth**
  - We need to double the size of table, add new bucket, redistribute data, update pointers
- If **local depth < global depth**
  - Add new bucket, redistribute data, update pointers

# Practice Question

**Q. What is the current hash function?**



# Practice Question



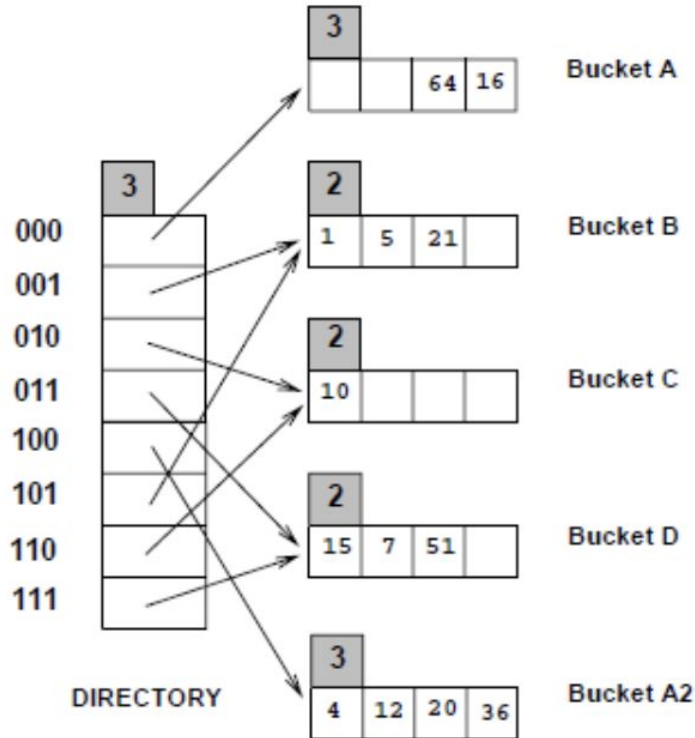
**Q. What is the current hash function?**

A.  $\text{hash}(x)$  = the most right three digits in binary representation of  $x$ .

Or


$\text{hash}(x) = x \bmod 8$

# Practice Question



**Q. What can you say about the last entry whose insertion into the index caused a split?**

# Practice Question




**Q. What can you say about the last entry whose insertion into the index caused a split?**

A. From the observation, we could conclude that the entry the last two digits of binary representation of which is '00' caused a split since the global depth is currently 3 and there are only two buckets the local depth of which are 3 corresponding to the entries the last two digits of the binary representation of which are 2. Thus the last entry whose insertion caused the split must be in {4, 12, 16, 20, 36, 64}.




# Practice Question



**Q. Which of the following insertion will cause  $|\text{global\_depth} - \text{local\_depth}|$  to be maximum among all possible cases?**

1. 13
2. 50
3. 55
4. 44


# Practice Question



**Q. Which of the following insertion will cause  $|\text{global\_depth} - \text{local\_depth}|$  to be maximum among all possible cases?**

1. 13
2. 50
3. 55
4. 44

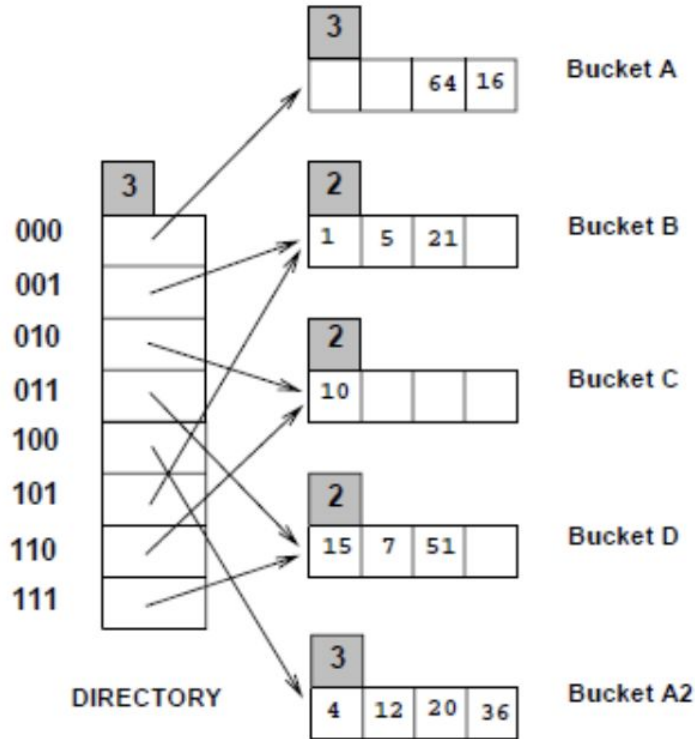
# Practice Question



**Q. Let  $L[i] = \text{local\_depth}[i]$ , where  $i$  denotes Bucket  $i$ . Let  $M = \max(L)$ . What is minimum possible size of the directory?**

1.  $2^{M-1}$
2.  $\log_2(M)$
3.  $\log_2(M-1)$
4.  $2^M$

# Practice Question



**Q. Let  $L[i] = \text{local\_depth}[i]$ , where  $i$  denotes Bucket  $i$ . Let  $M = \max(L)$ . What is minimum possible size of the directory?**

1.  $2^{M-1}$
2.  $\log_2(M)$
3.  $\log_2(M-1)$
4.  $2^M$

# Practice Question

**Q. True/False: A Bucket will have more than one pointers pointing to it if its local depth is less than the global depth.**

**True**

**Q. True/False: If Local Depth of the overflowing bucket is equal to the global depth, only then the directories are doubled and the global depth is incremented by 1.**

**True**

# Limitations

- The directory size may increase significantly if several records are hashed on the same bucket.
- Size of every bucket is fixed.
- Memory is wasted in pointers when the global depth and local depth difference becomes drastic.

# External Sorting

- why sorting?
  - bulk loading in B+ tree, duplicate elimination, users want the data sorted, useful in join algorithms e.t.c
- what about the standard algorithms?
  - merge sort, quick sort, heap sort → in-memory
  - the data typically does not fit in memory
  - e.g how to sort 1TB data with 8GB of RAM?

# External Merge

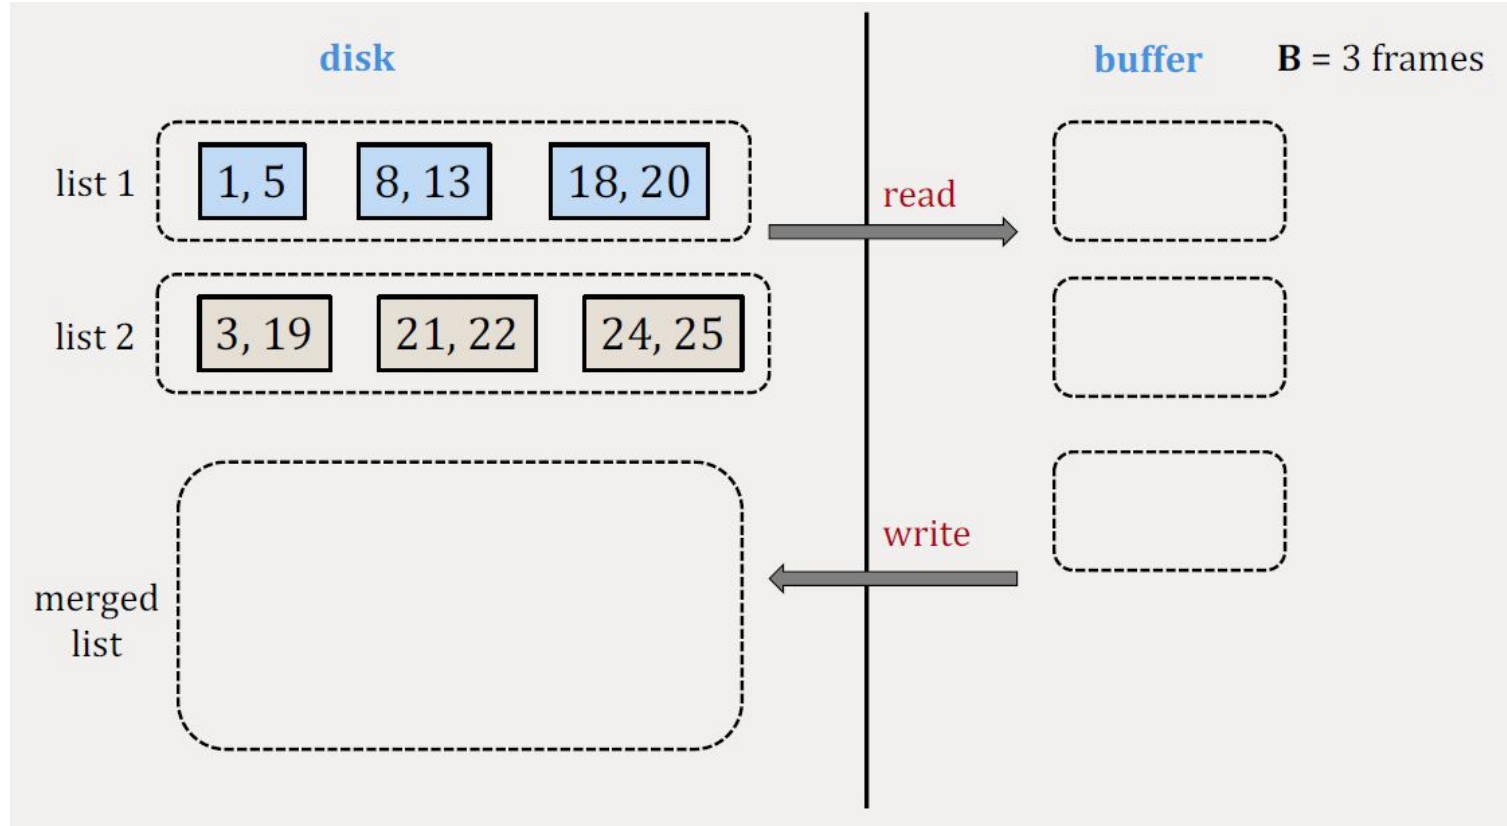
- **Input:** 2 sorted lists (with  $M$  and  $N$  pages)
- **Output:** 1 merged sorted list ( with  $M+N$  pages)

**Can we do the merging efficiently in terms of I/O using a buffer pool of size at least 3?**

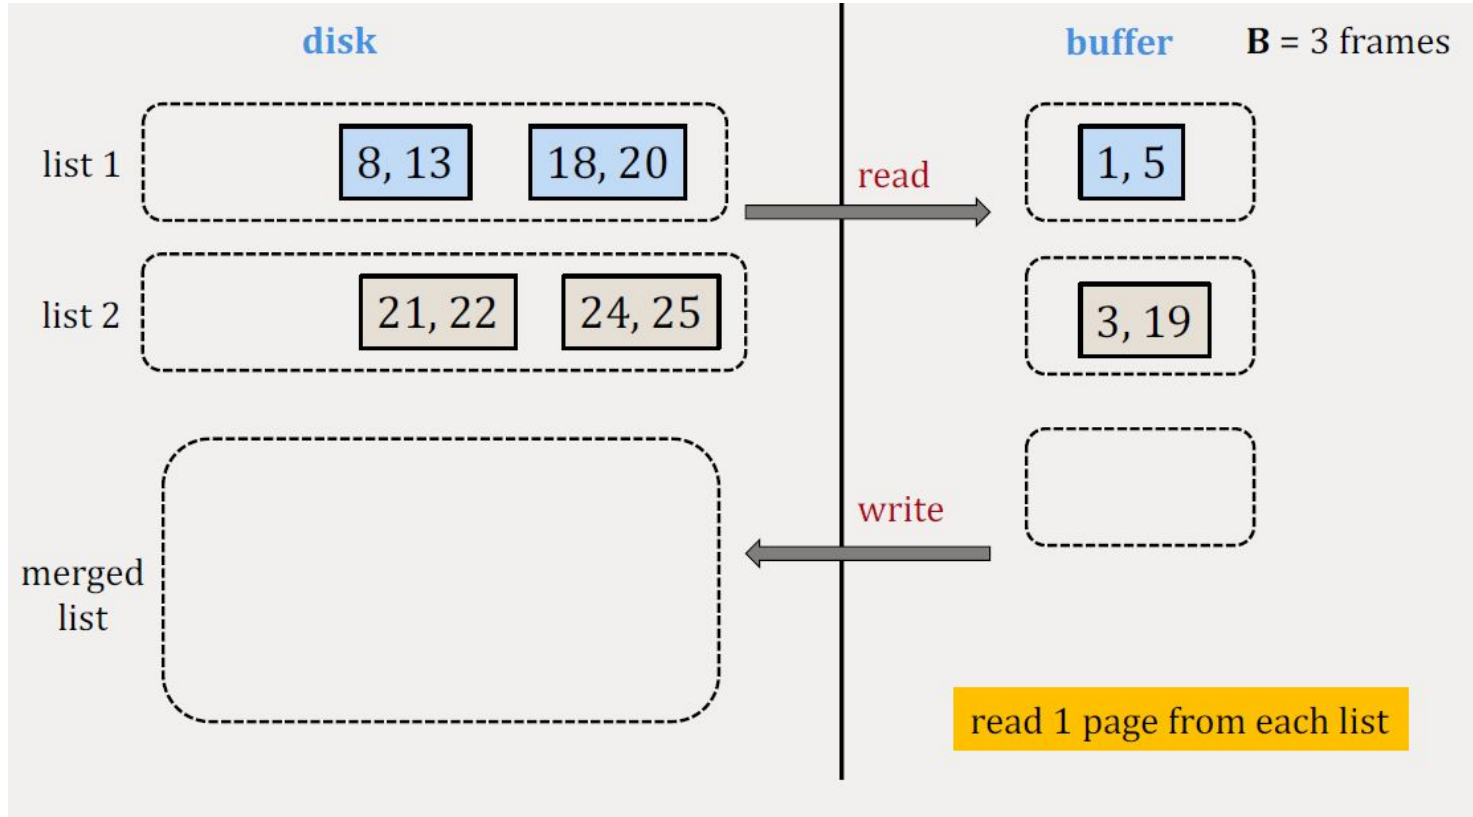
- **Yes, using only  $2(M+N)$  I/Os!**



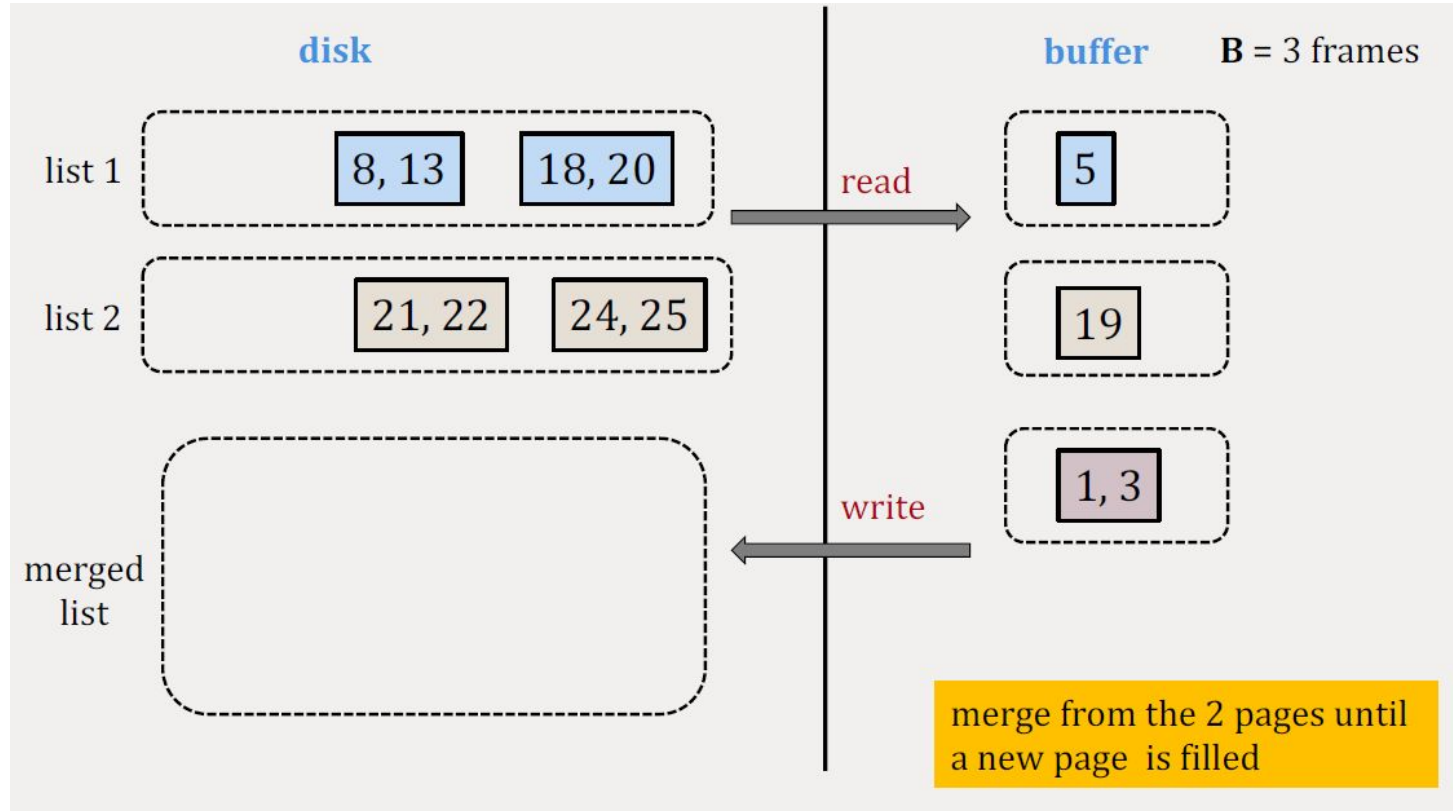
# External Merge Example(1)



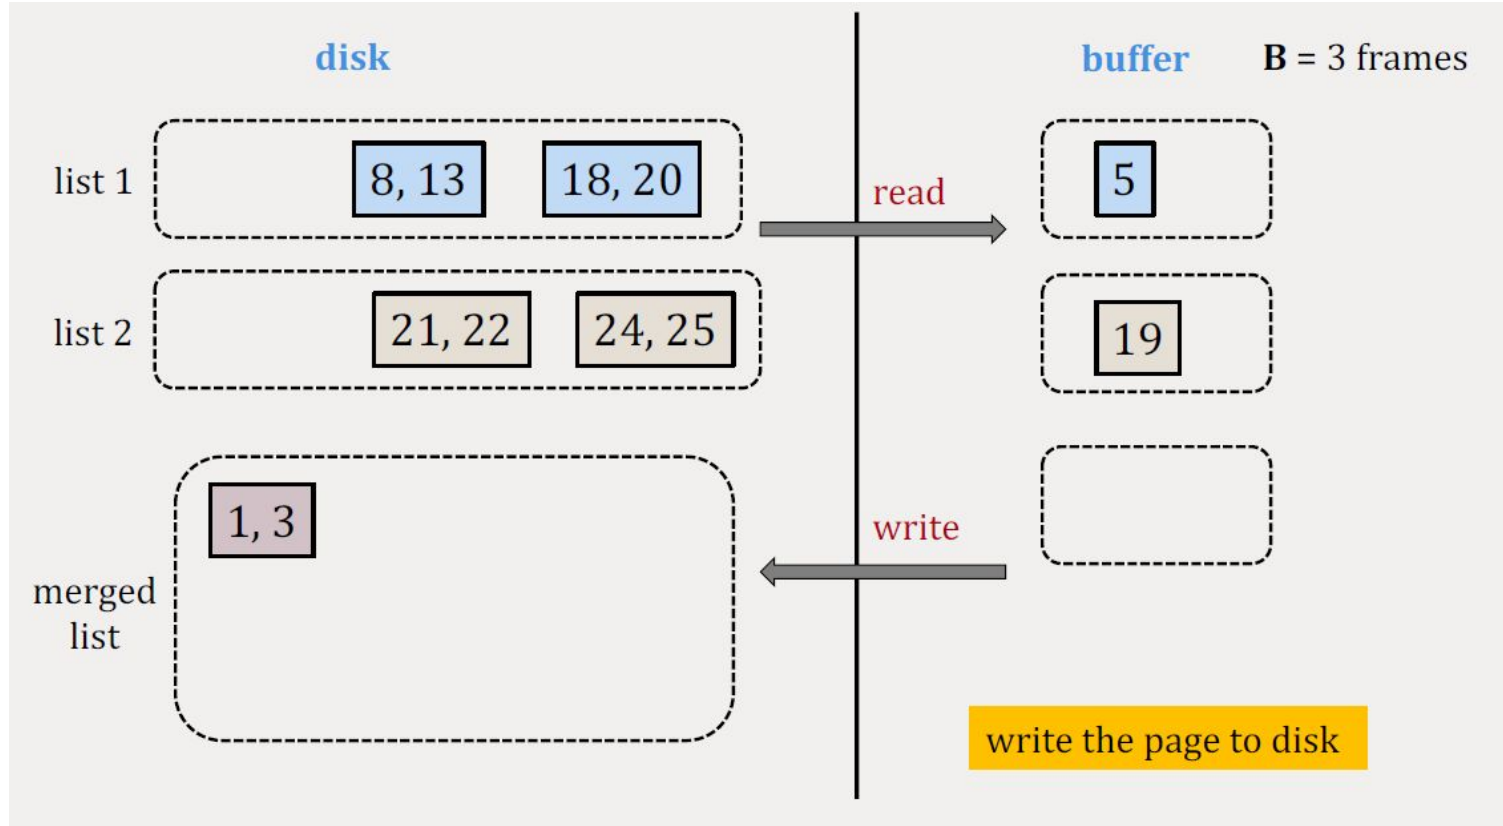
# External Merge Example(2)



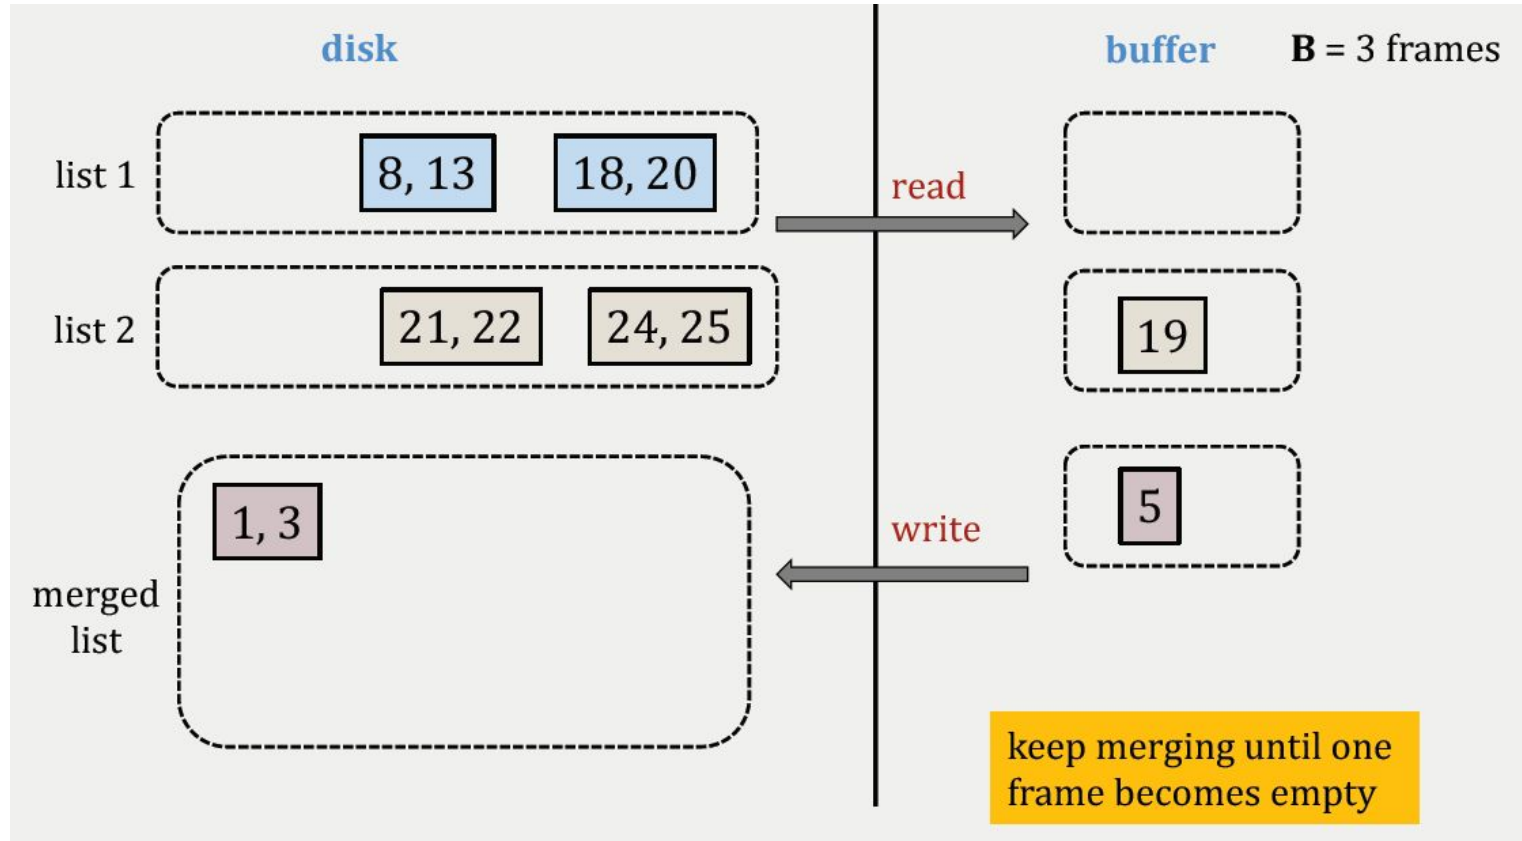
# External Merge Example(3)



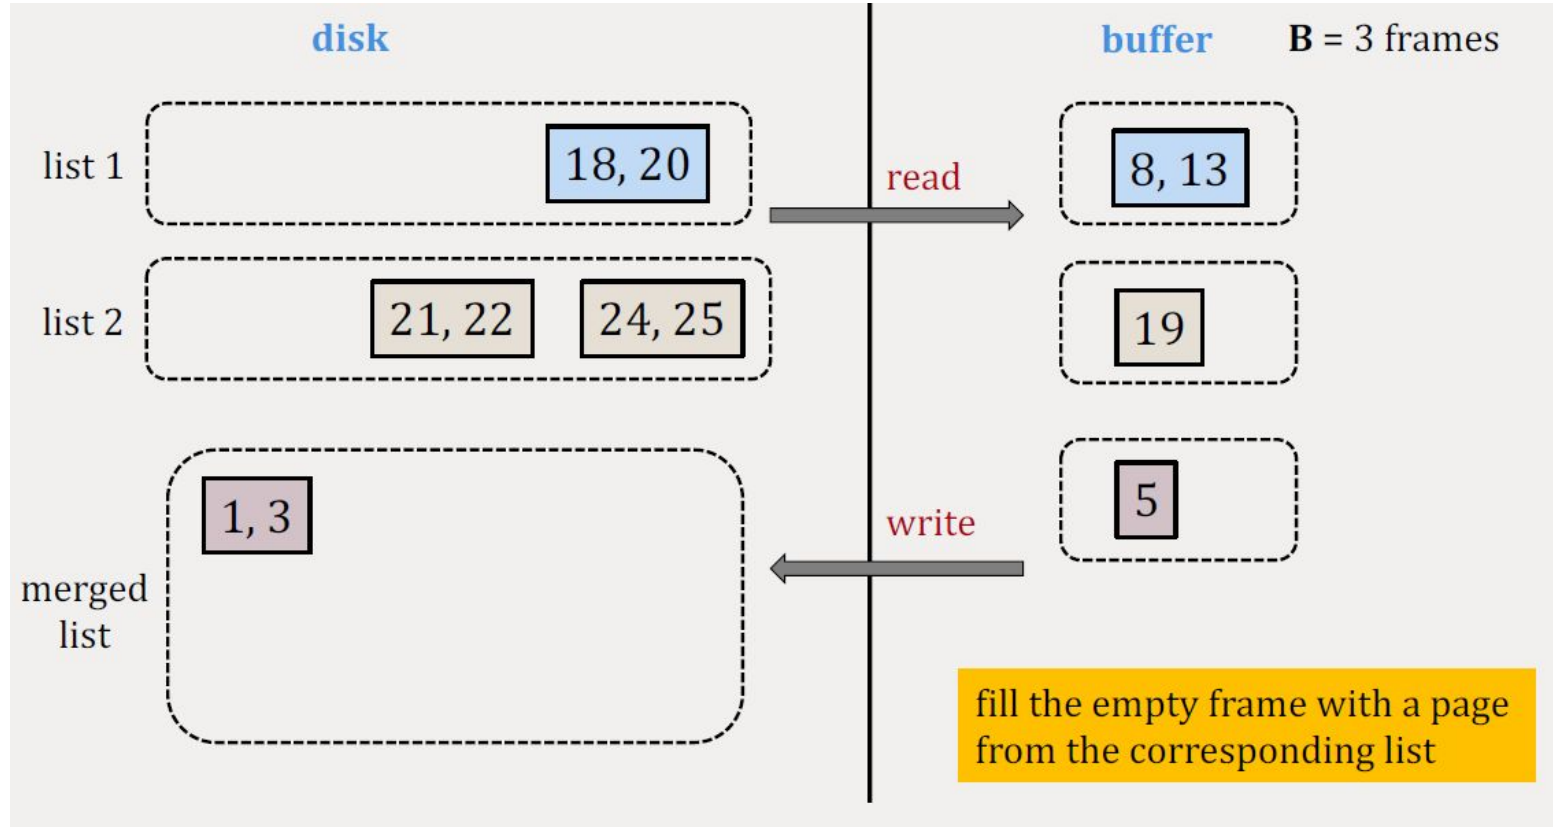
# External Merge Example(4)



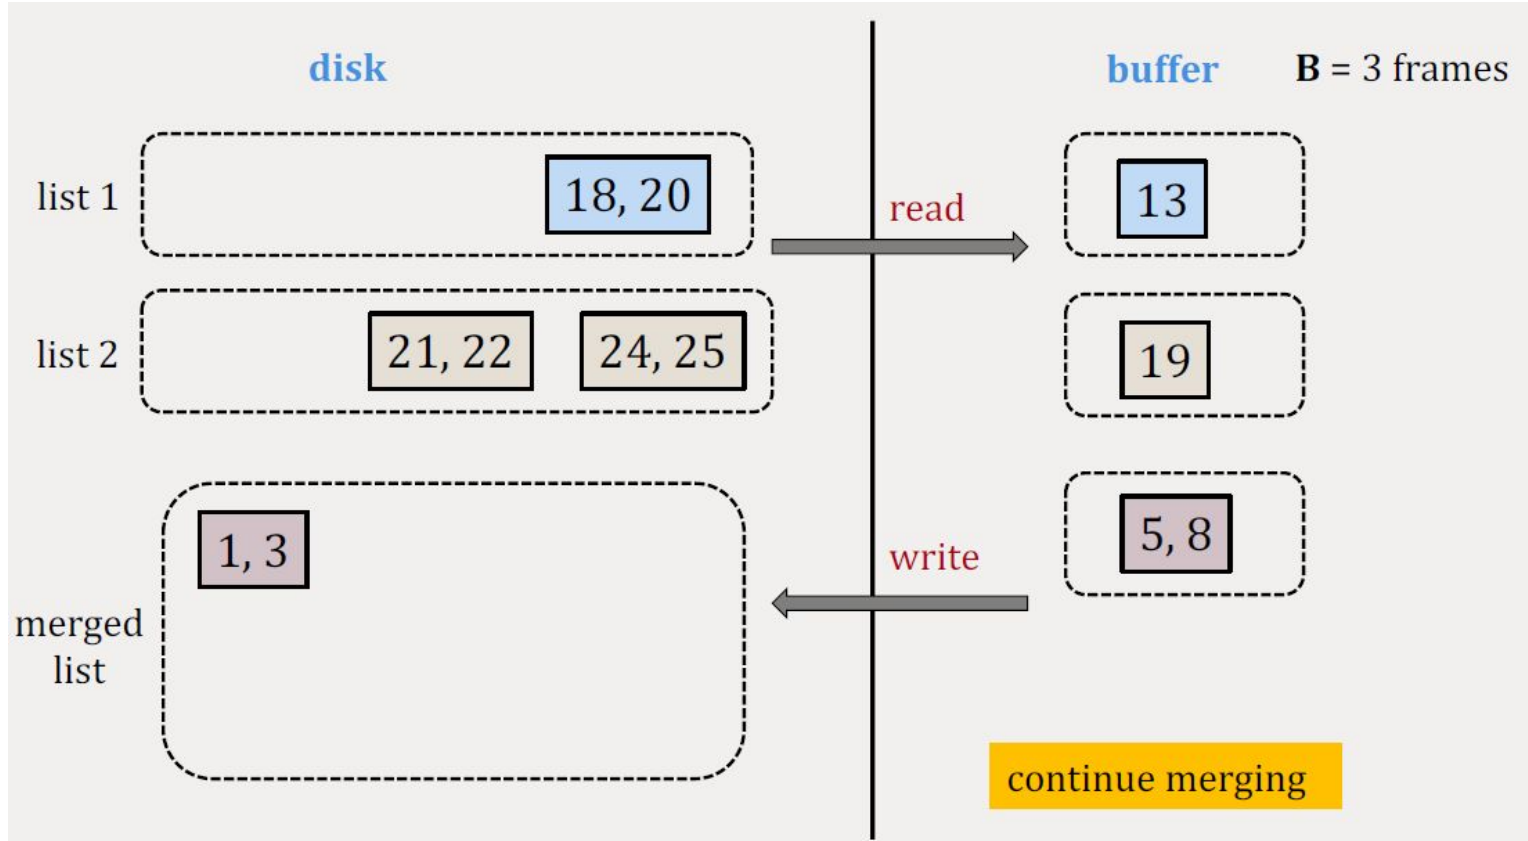
# External Merge Example(5)



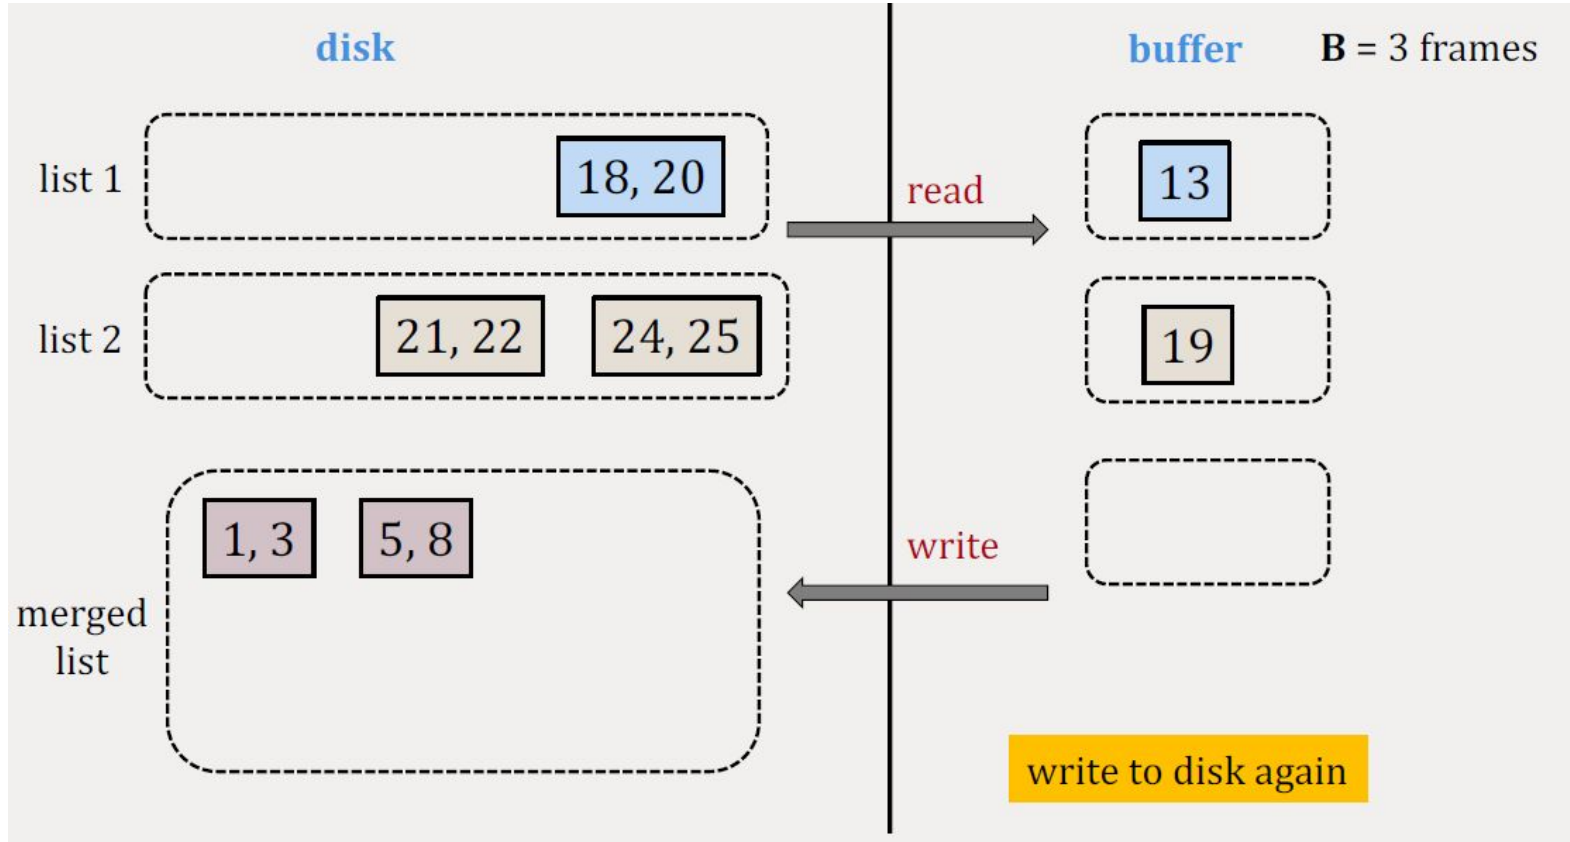
# External Merge Example(6)



# External Merge Example(7)

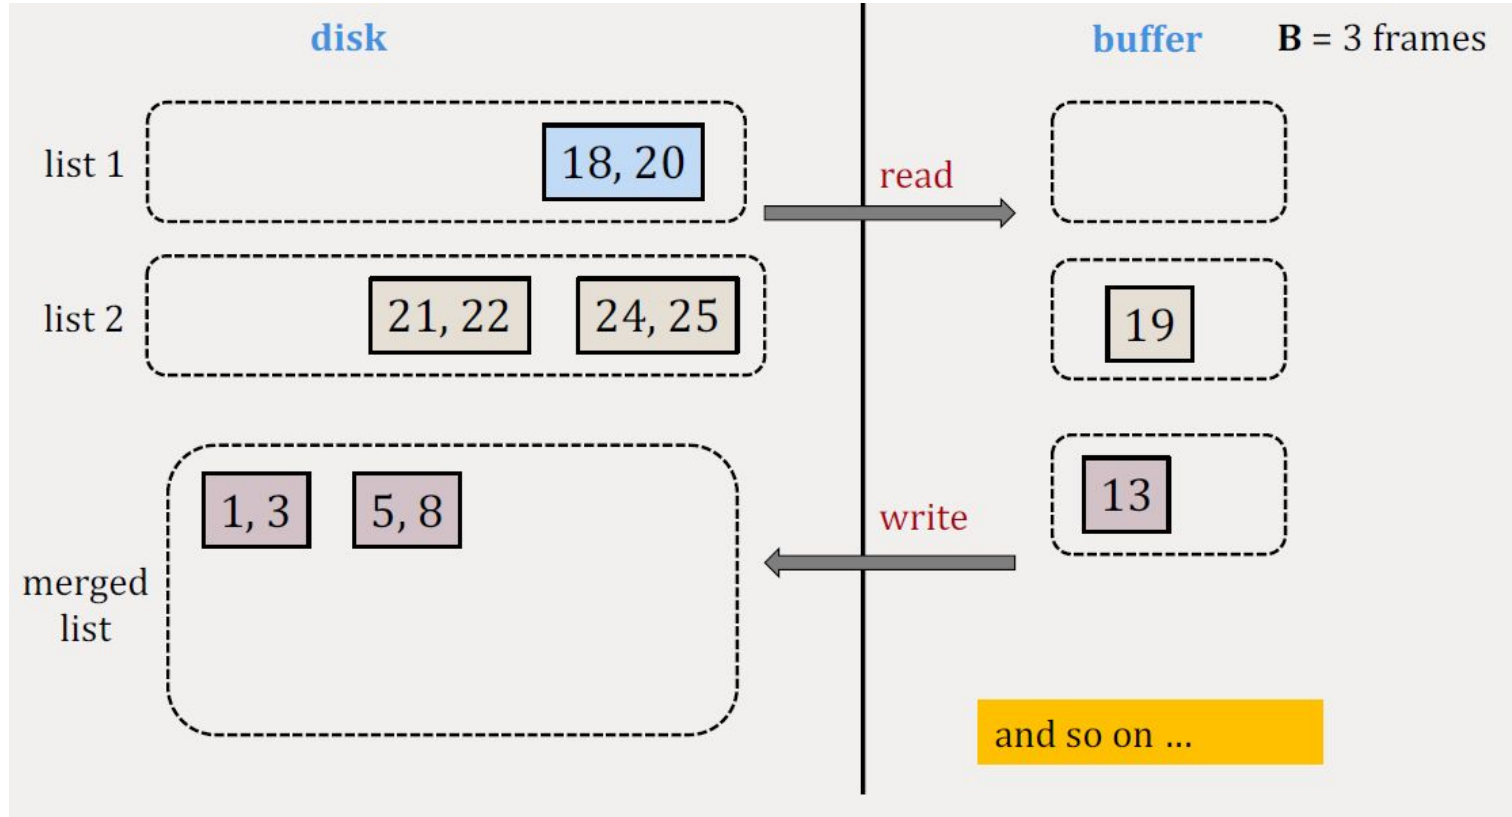


# External Merge Example(8)





# External Merge Example(9)



# The Sorting Problem

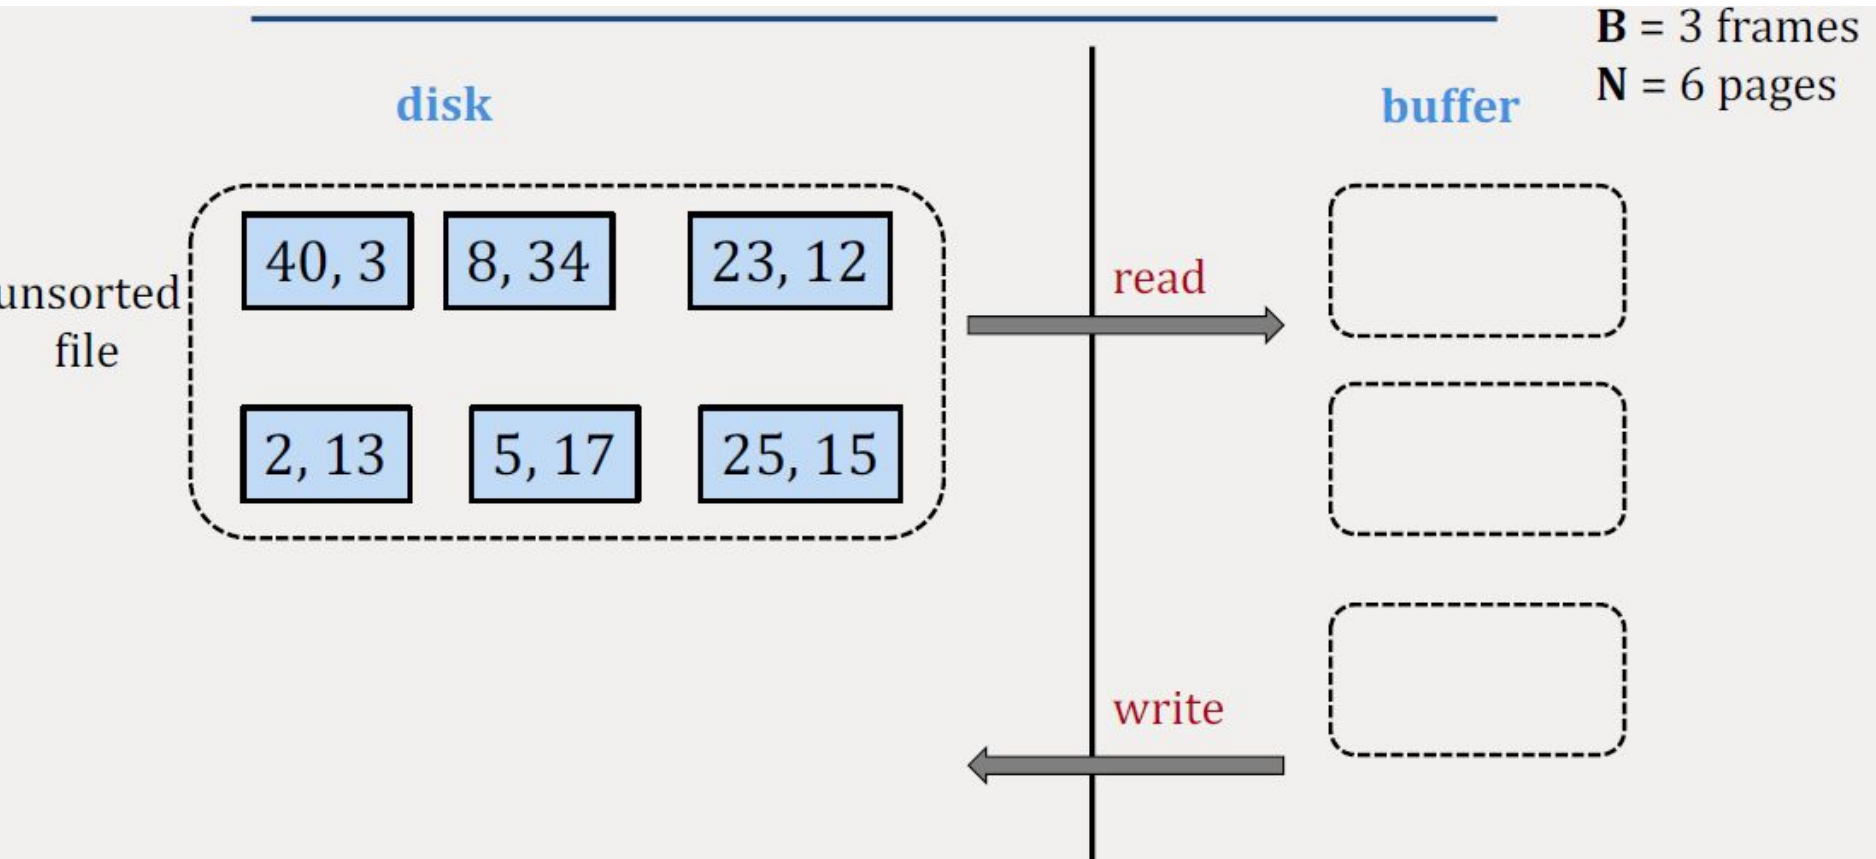
- **B** available pages in the buffer pool
- a relation  $R$  of **N** pages (where  $N > B$ )

**SORTING**: output the same relation sorted on a given attribute

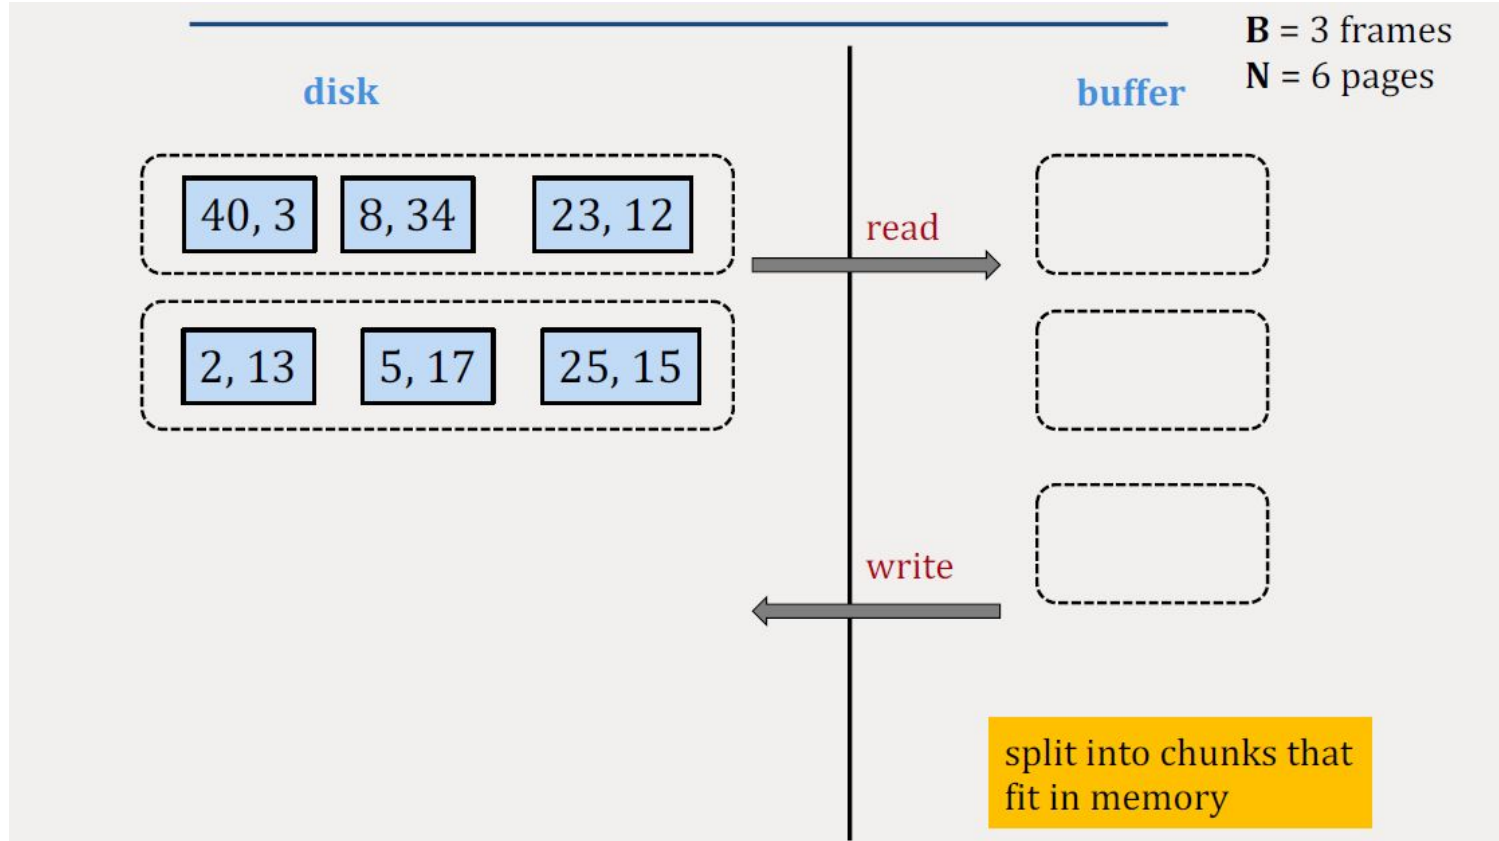
# Key Idea

- split into chunks small enough to sort in memory (called runs)
- merge groups of runs using the **external merge algorithm**
- keep merging the resulting runs (each time is called a pass) until left with a single sorted file

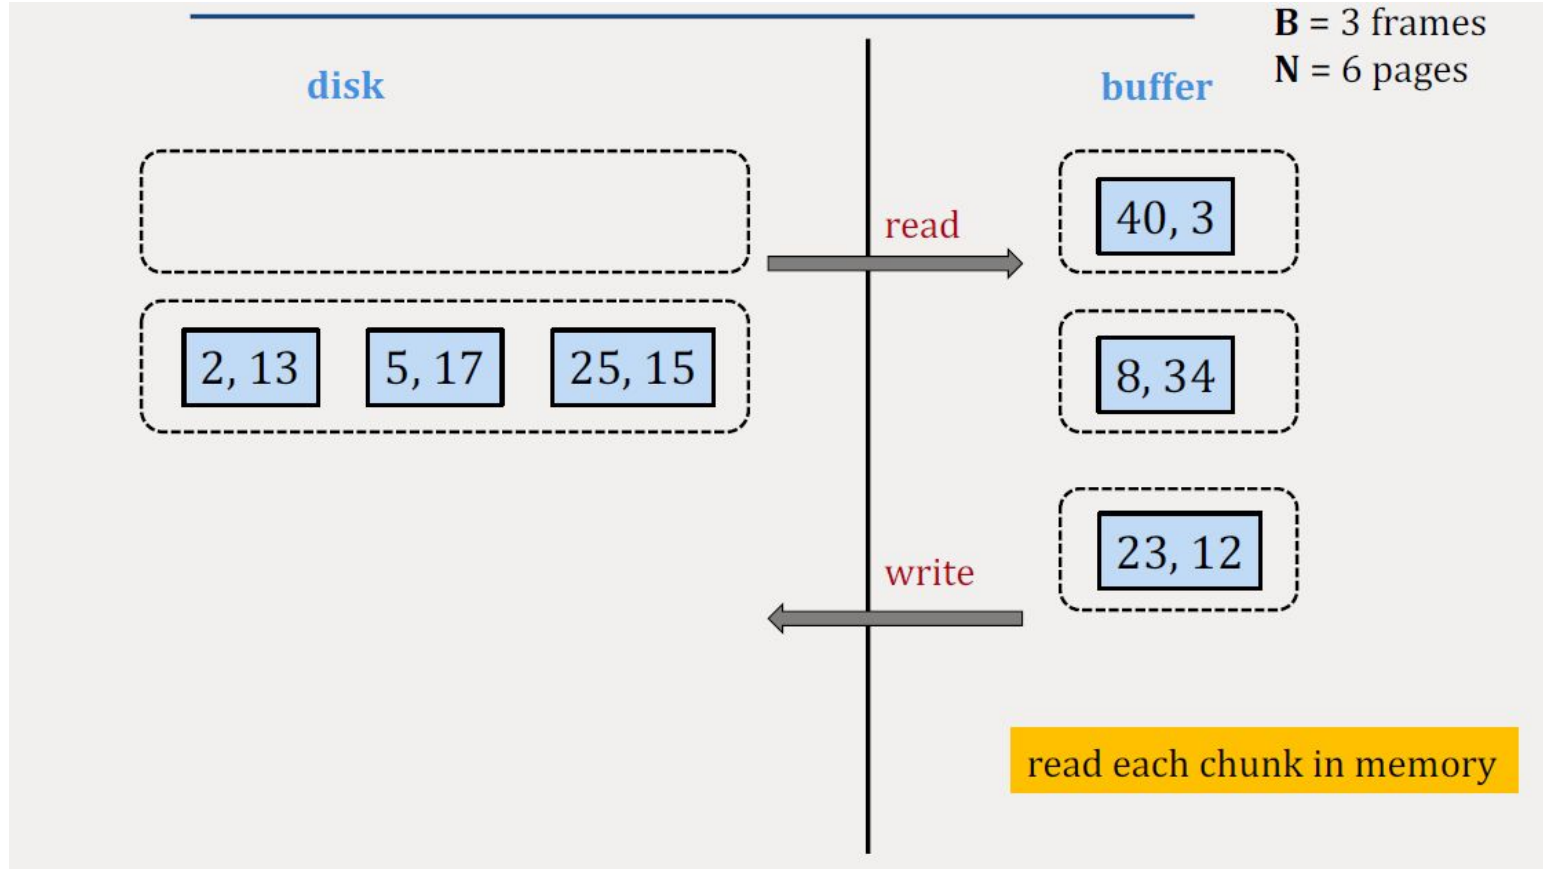
## 2-Way Sort Example(1)



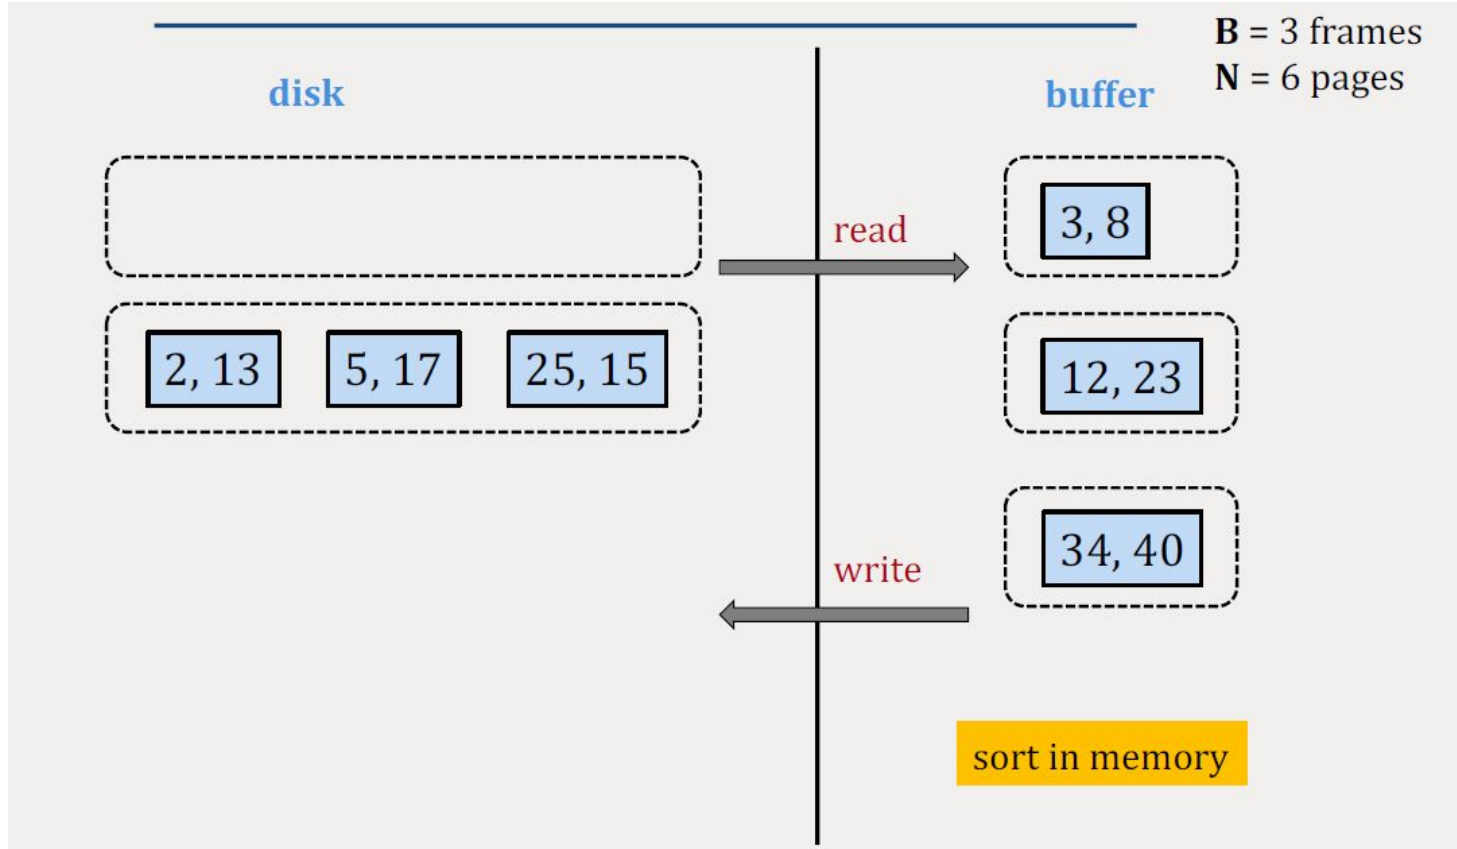
## 2-Way Sort Example(2)



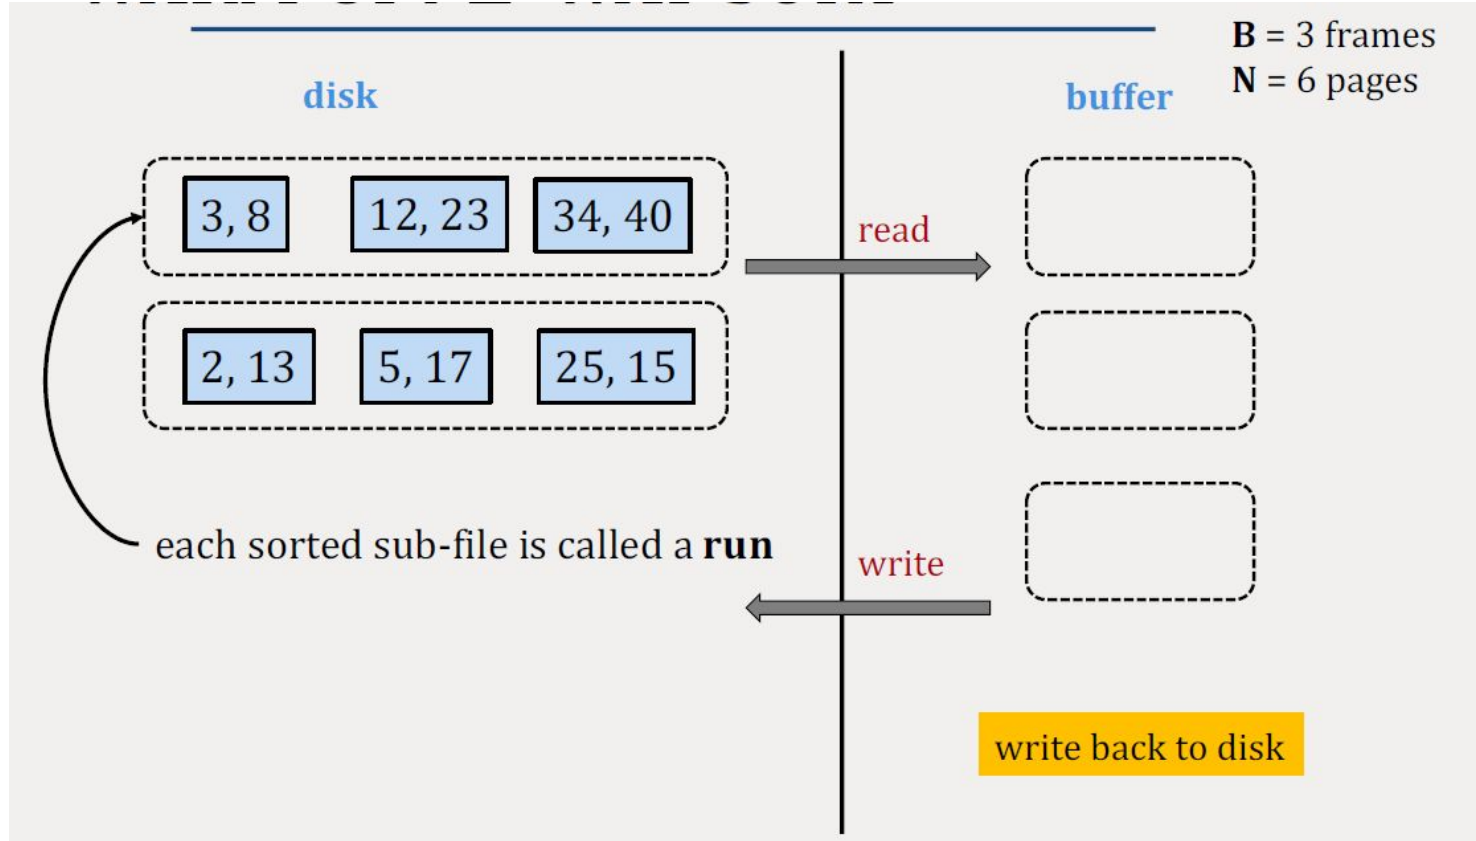
## 2-Way Sort Example(3)



## 2-Way Sort Example(4)

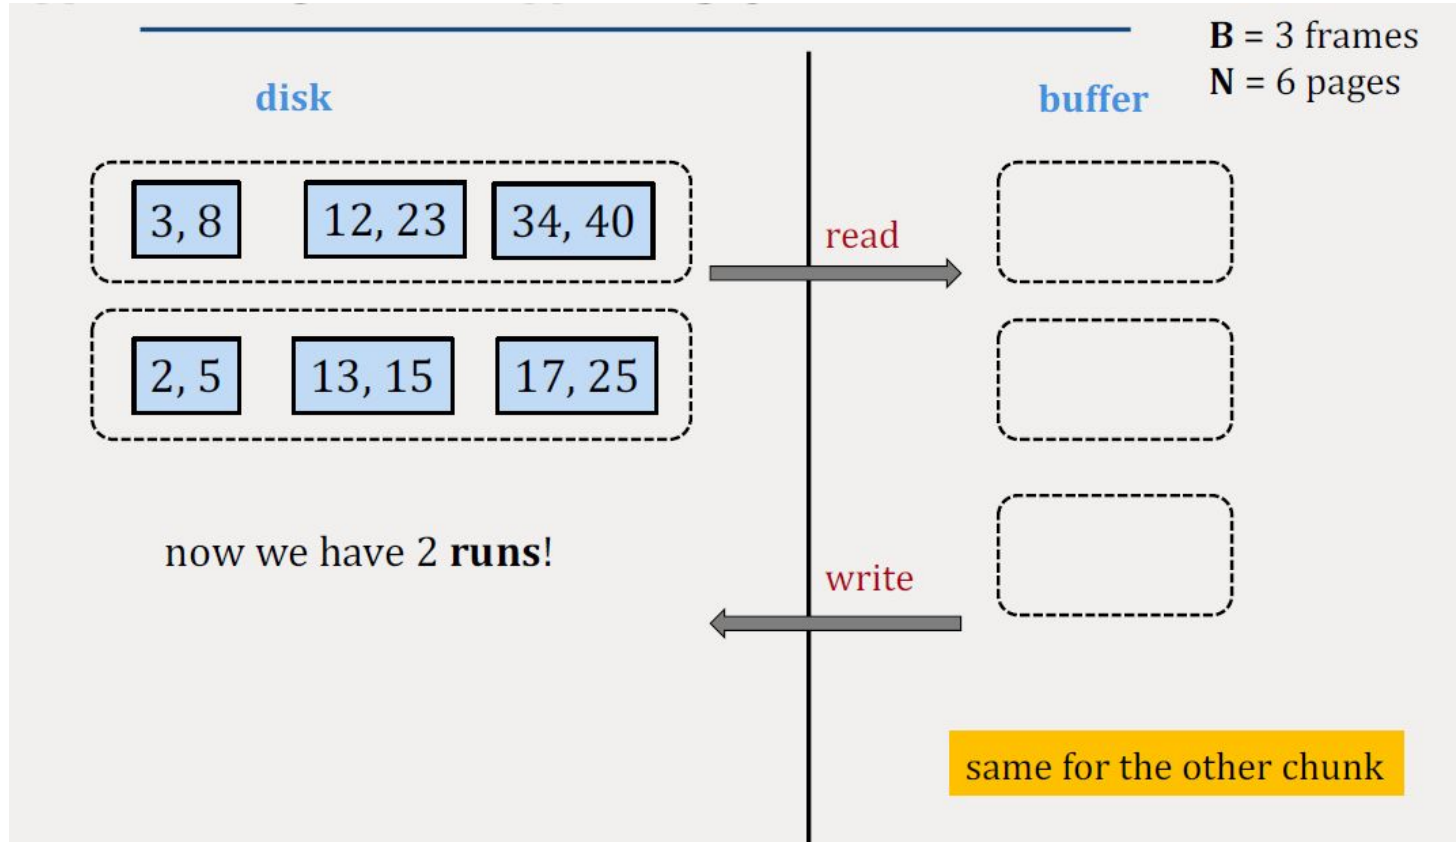


## 2-Way Sort Example(5)

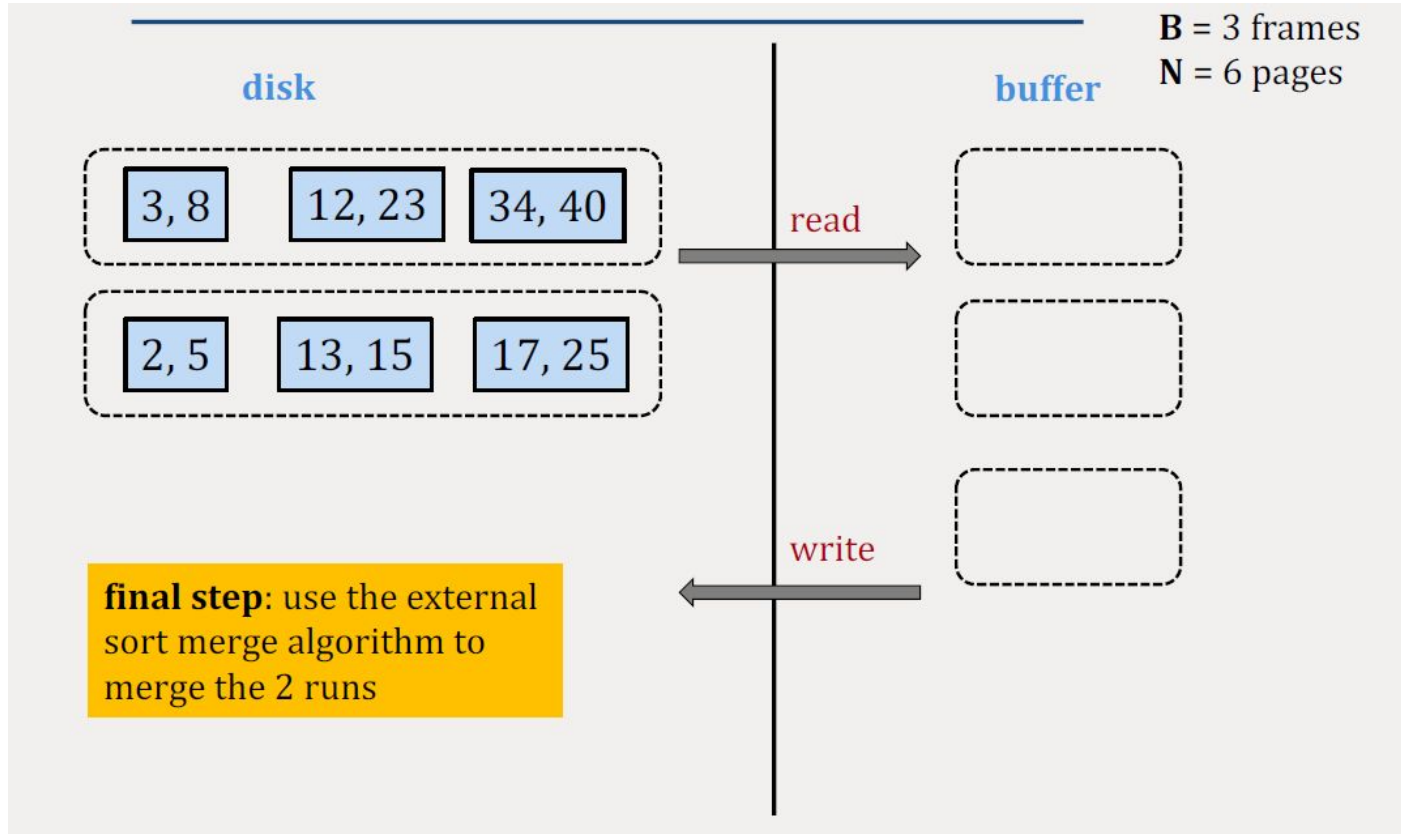




## 2-Way Sort Example(6)



## 2-Way Sort Example(7)



# Calculating the I/O cost

In our example, **B**=3 (buffer pool pages), **N**=6 pages

- **Pass 0: sorting**
  - 1 read + 1 write for every page
  - total cost =  $6 * (1+1) = 12$  I/Os
- **Pass 1: merging**
  - total cost =  $2 * (3+3) = 12$  I/Os

Total: 24 I/Os

# External Sorting Problem

Given two files A and B and each file includes 4 pages (each page contains 4 entries only) as shown as the following:

A: [[1,4,2, 4], [5,2,18,9], [28,7,16,5], [13, 14, 11, 12]]

B: [[9,5,2,1], [25,5, 31, 0], [45, 33, 12, 16], [19, 13, 16, 21]]

Assume you have 3 buffer pages in main memory, show the “external sort” procedure for the given example.

# Solution (1)

- **First Run (Sort Phase)** → split into chunks of 3 (buffer pages 3), bring each chunk into memory and apply the in-memory sorting algorithm, write it back to disk. End up with 3 sorted chunks in disk.

Input:

A: [1,4,2, 4], [5,2,18,9], [28,7,16,5], [13, 14, 11, 12]]

B: [9,5,2,1], [25,5, 31, 0], [45, 33, 12, 16], [19, 13, 16, 21]]

Output:

[1, 2, 2, 4] [4, 5, 5, 7] [9, 16, 18, 28]

[0, 1, 2, 5] [5, 9, 11, 12] [13, 14, 25, 31]

[12, 13, 16, 16] [19, 21, 33, 45]

**I/O cost:** read & write of 8 pages → 16 I/O  
(in terms of pages)

# Solution (2)

- **Second Run (Merge Phase)** → Take two pages each time from the sorted chunks into memory, merge them and write them to disk

Input:

[1, 2, 2, 4] [4, 5, 5, 7] [9, 16, 18, 28]

[0, 1, 2, 5] [5, 9, 11, 12] [13, 14, 25, 31]

[12, 13, 16, 16] [19, 21, 33, 45]

Output:

[0, 1, 1, 2] [2, 2, 4, 4] [5, 5, 5, 5] [7, 9, 9, 11] [12, 13, 14, 16] [18, 25, 28, 31]

[12, 13, 16, 16] [19, 21, 33, 45]

**I/O cost:** read & write of 6 pages => 12 I/O

# Solution (3)

- **Third Run (Merge Phase)** → Take two pages each time from the sorted chunks into memory, merge them and write them to disk

Input:

[0, 1, 1, 2] [2, 2, 4, 4] [5, 5, 5, 5] [7, 9, 9, 11] [12, 13, 14, 16] [18, 25, 28, 31]

[12, 13, 16, 16] [19, 21, 33, 45]

Output:

[0, 1, 1, 2] [2, 2, 4, 4] [5, 5, 5, 5] [7, 9, 9, 11] [12, 12, 13, 13] [14, 16, 16, 16] [18, 19, 21, 25] [28, 31, 33, 45]

**I/O cost:** read & write of 6 pages => 16 I/O

**Total I/O cost: 44 I/O**

Thanks!